

COMP 6461

Computer Networks & Protocols

Winter 2023

Dr. Abdelhak Bentaleb



Lecture 3a

Application Layer (Part 2)

Application layer: overview

- Principles of network applications
- video streaming and content distribution networks
- **Web and HTTP**
- E-mail, SMTP, IMAP
- The Domain Name System DNS
- P2P applications
- socket programming with UDP and TCP

Web and HTTP

First, a quick review...


- web page consists of *objects*, each of which can be stored on different Web servers
- object can be HTML file, JPEG image, Java applet, audio file,...
- web page consists of *base HTML-file* which includes *several referenced objects, each* addressable by a *URL*, e.g.,

www.someschool.edu / someDept/pic.gif
host name path name

HTTP overview

HTTP: HyperText Transfer Protocol

- Web application's layer protocol
- Client/Server model
- **Client:** browser that requests, receives (using HTTP protocol), and displays web objects.
- **Server:** web server that sends (using HTTP protocol) objects in response to requests.



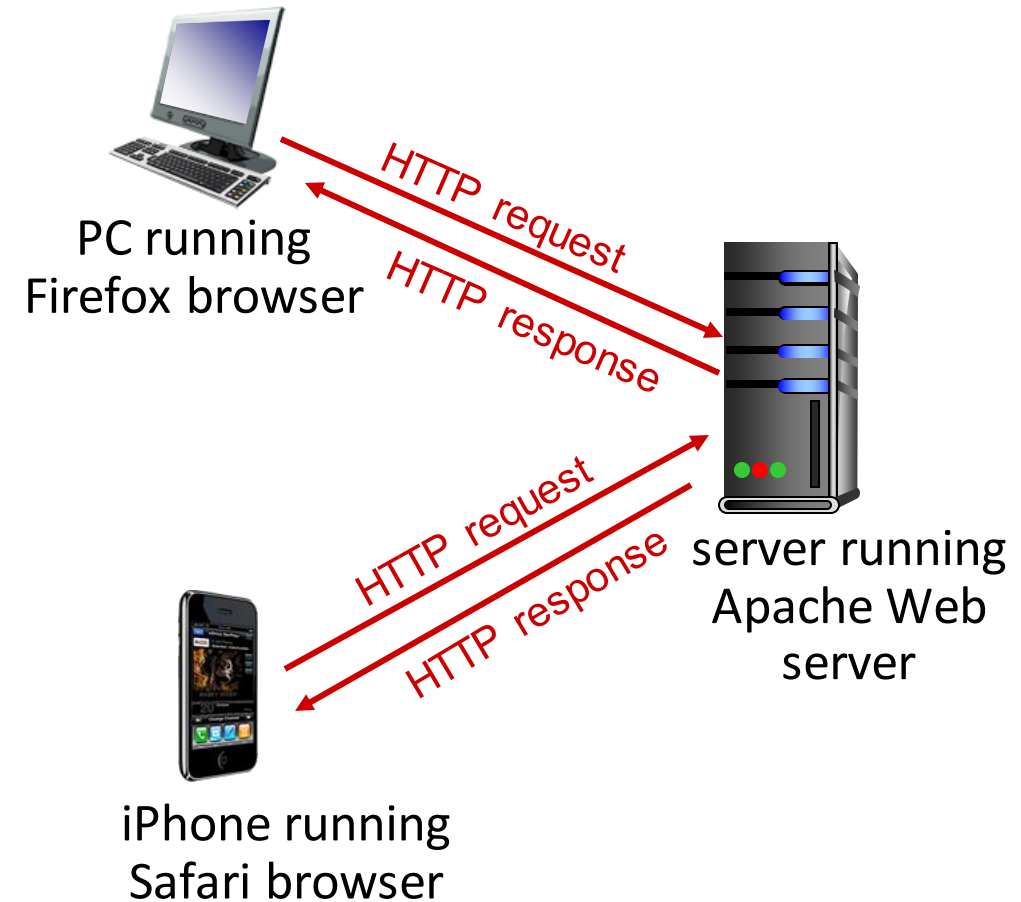
HTTP overview (continued)

HTTP uses TCP (at transport layer)

- client initiates TCP connection (create socket) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application layer protocol messages) exchanges between browser (HTTP client) and web server (HTTP server)
- Close TCP connection

HTTP is “stateless”

- server maintains *no* information about past client requests



HTTP connections: two types

Non-persistent HTTP

1. TCP connection opened
2. at most one object sent over TCP connection
3. TCP connection closed

downloading multiple objects required multiple connections

Persistent HTTP

- TCP connection opened to a server
- multiple objects can be sent over *single* TCP connection between client, and that server
- TCP connection closed

Non-persistent HTTP: example

User enters URL: `www.someSchool.edu/someDepartment/home.index`
(containing text, references to 10 jpeg images)



1a. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80



1b. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80 “accepts” connection, notifying client

2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`

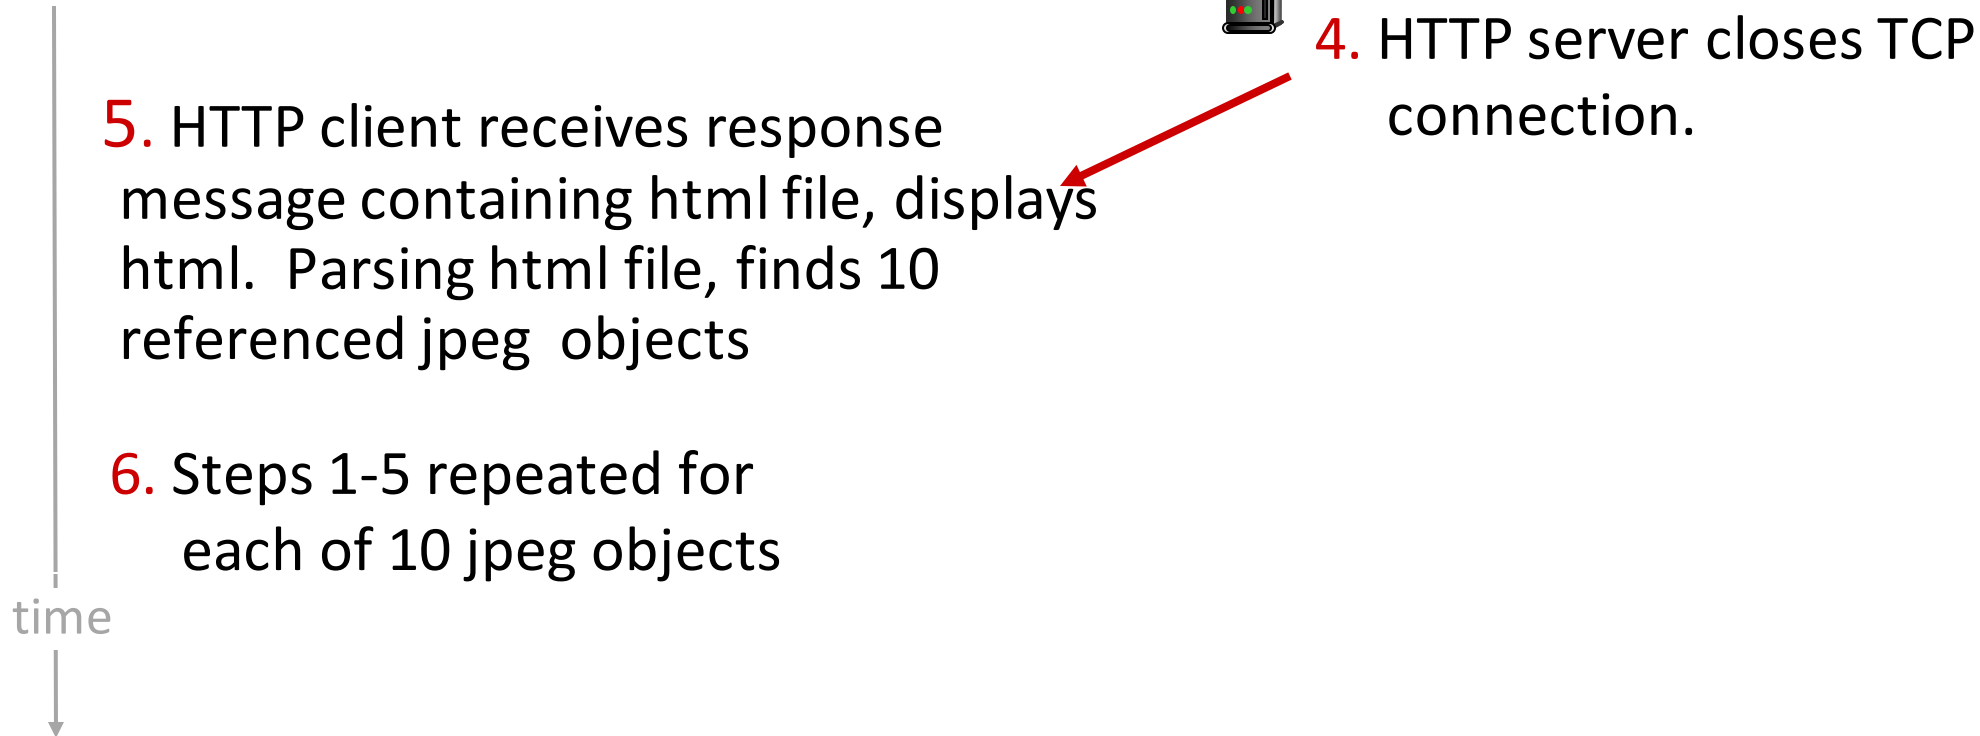
3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket

time



Non-persistent HTTP: example (cont.)

User enters URL: `www.someSchool.edu/someDepartment/home.index`
(containing text, references to 10 jpeg images)

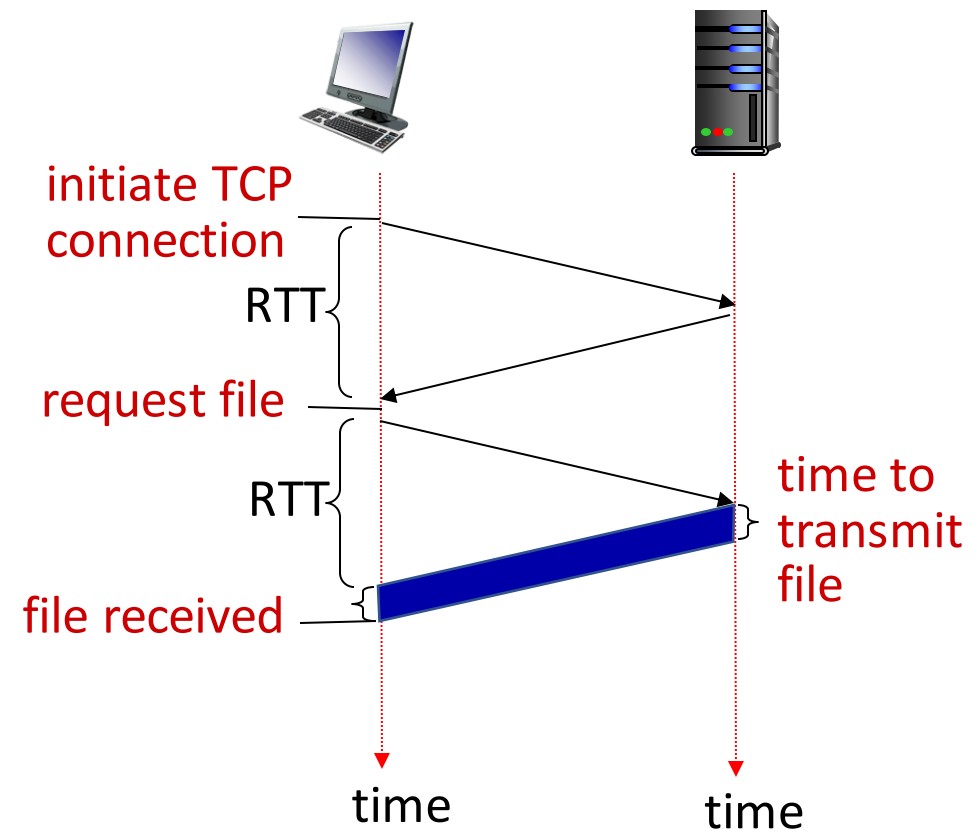


Non-persistent HTTP: response time

RTT (definition): time for a small packet to travel from client to server and back

HTTP response time (per object):

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- object/file transmission time



Non-persistent HTTP response time = $2RTT + \text{file transmission time}$

Persistent HTTP (HTTP 1.1)

Non-persistent HTTP issues:


- requires 2 RTTs per object
- OS overhead for *each* TCP connection
- browsers often open multiple parallel TCP connections to fetch referenced objects in parallel



Persistent HTTP (HTTP1.1):


- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects (cutting response time in half)

HTTP request message

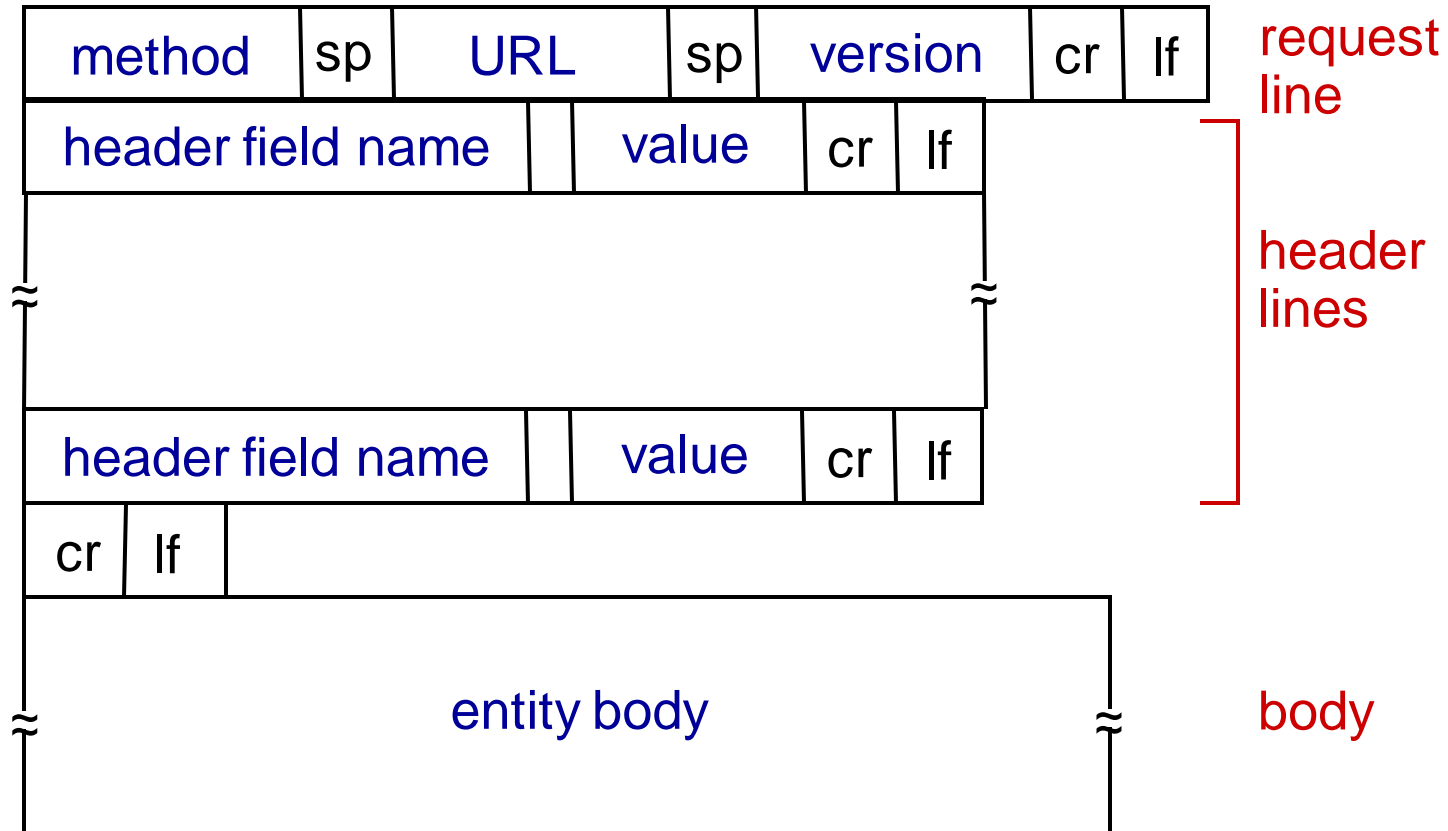
- two types of HTTP messages: *request, response*
- **HTTP request message:**
 - ASCII (human-readable format)

request line (GET, 
POST,
HEAD commands)

 carriage return character
 line-feed character

carriage return, line feed 
at start of line indicates
end of header lines

HTTP request message: general format



Other HTTP request messages

POST method:

- web page often includes form input
- user input sent from client to server in entity body of HTTP POST request message

GET method (for sending data to server):

- include user data in URL field of HTTP GET request message (following a '?'):

`www.somesite.com/animalsearch?monkeys&banana`

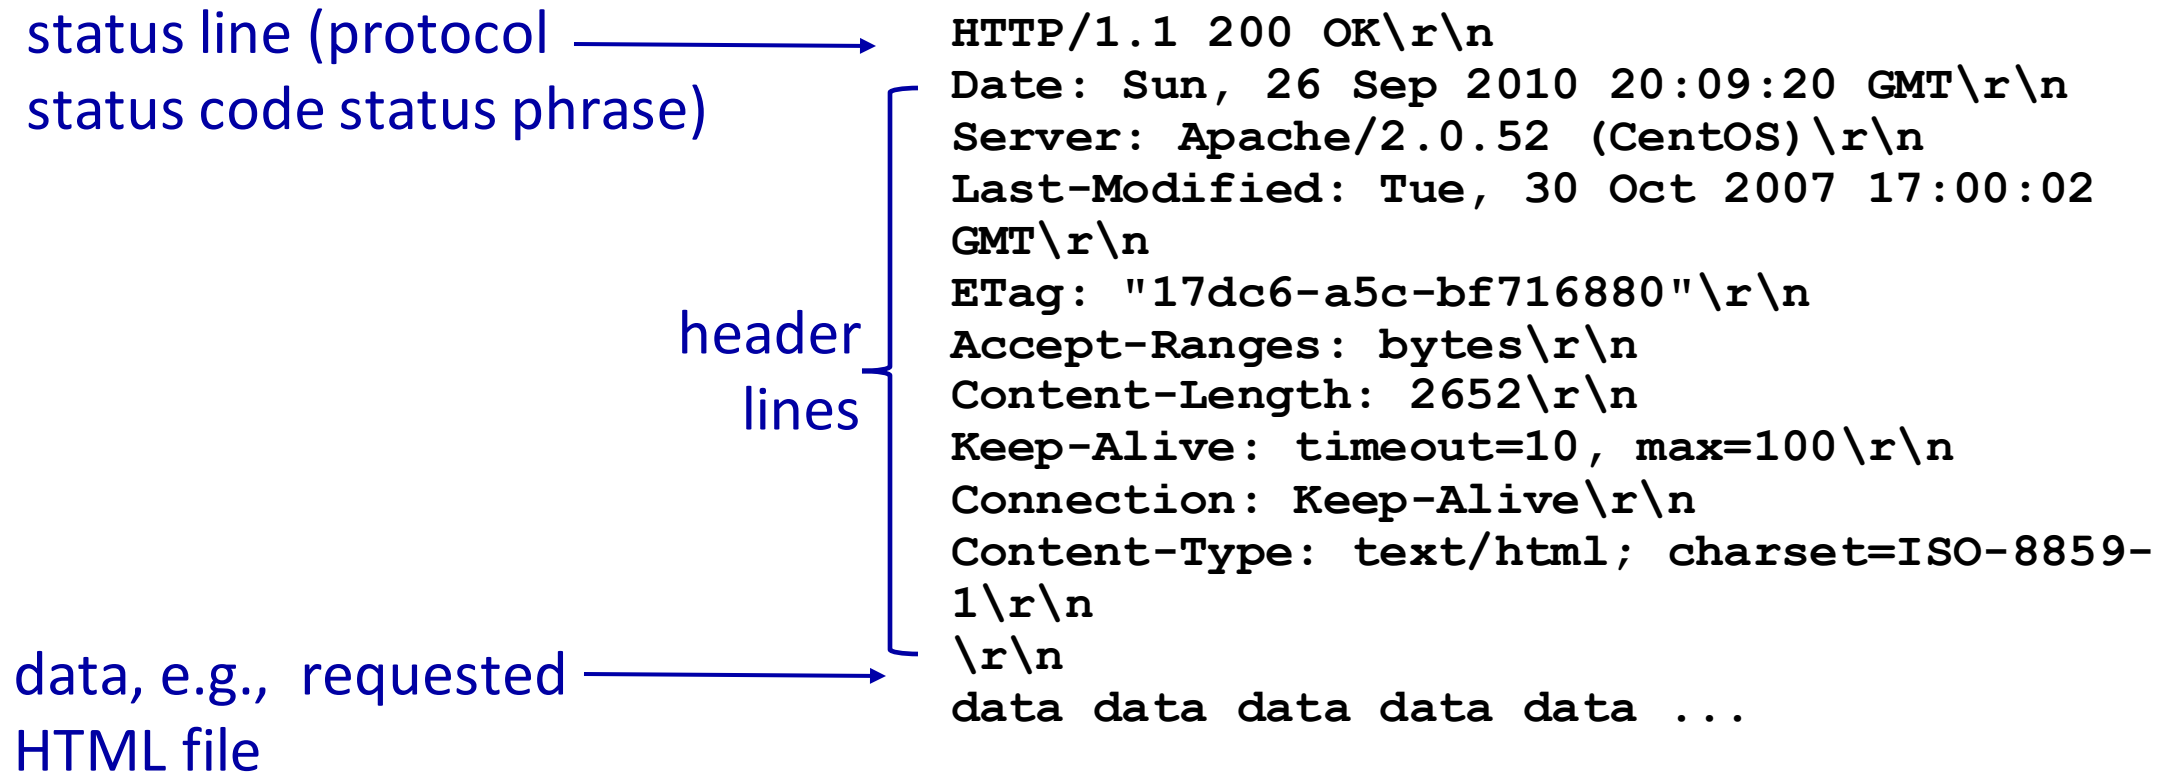
HEAD method:

- requests headers (only) that would be returned *if* specified URL were requested with an HTTP GET method.

PUT method:

- uploads new file (object) to server
- completely replaces file that exists at specified URL with content in entity body of POST HTTP request message

HTTP response message



* Check out the online interactive exercises for more examples: http://gaia.cs.umass.edu/kurose_ross/interactive/

HTTP response status codes

- status code appears in 1st line in server-to-client response message.
- some sample codes:

200 OK

- request succeeded, requested object later in this message

301 Moved Permanently

- requested object moved, new location specified later in this message (in Location: field)

400 Bad Request

- request msg not understood by server

404 Not Found

- requested document not found on this server

505 HTTP Version Not Supported

Maintaining user/server state: cookies

Recall: HTTP GET/response interaction is *stateless*

Web sites and client browser use *cookies* to maintain some state between transactions

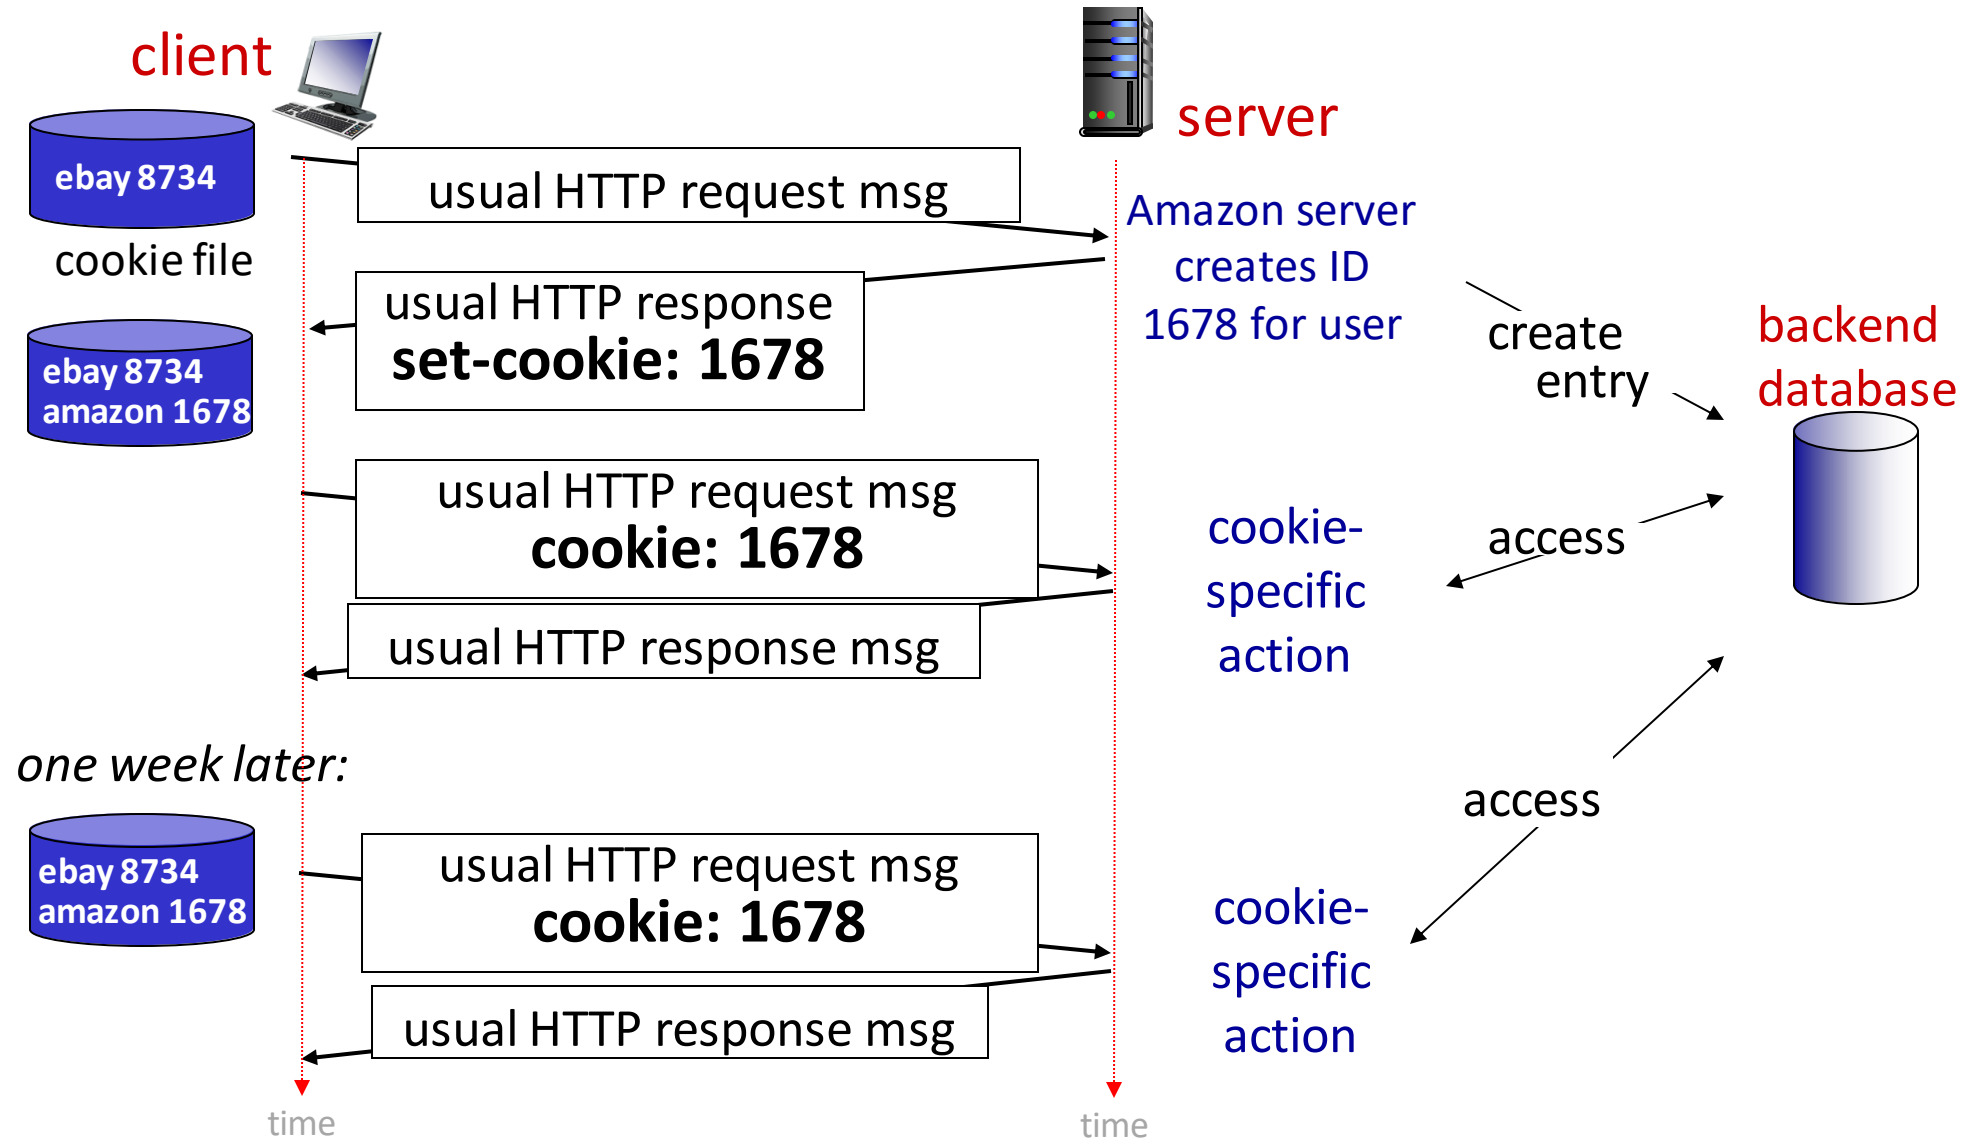
four components:

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in next HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

Example:

- Susan uses browser on laptop, visits specific e-commerce site for first time
- when initial HTTP requests arrives at site, site creates:
 - unique ID (aka “cookie”)
 - entry in backend database for ID
- subsequent HTTP requests from Susan to this site will contain cookie ID value, allowing site to “identify” Susan


Maintaining user/server state: cookies



Web caches

Goal: satisfy client requests without involving origin server

- user configures browser to point to a (local) *Web cache*
- browser sends all HTTP requests to cache
 - *if* object in cache: cache returns object to client
 - *else* cache requests object from origin server, caches received object, then returns object to client



Web caches (aka proxy servers)

- Web cache acts as both client and server
 - server for original requesting client
 - client to origin server
- server tells cache about object's allowable caching in response header:

```
Cache-Control: max-age=<seconds>
```

```
Cache-Control: no-cache
```

Why Web caching?

- reduce response time for client request
 - cache is closer to client
- reduce traffic on an institution's access link

Caching example


Scenario:

- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits (0.1 Mbits)
- average request rate from browsers to origin servers: 15 req/sec
 - avg data rate to browsers: 1.50 Mbps (15 x 0.1)

Performance:

- access link utilization = **.97**
- LAN utilization: .0015 (15/10000)
- end-end delay = Internet delay +
access link delay + LAN delay
= 2 sec + **minutes** + usecs

*problem: large
queueing delays
at high utilization!*



Option 1: buy a faster access link


Scenario:

- access link rate: ~~1.54~~ 154 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
 - avg data rate to browsers: 1.50 Mbps

Performance:

- access link utilization = ~~.97~~ .0097
- LAN utilization: .0015
- end-end delay = Internet delay +
access link delay + LAN delay
= 2 sec + ~~minutes~~ + usecs

Cost: faster access link (expensive!) → msecs



Option 2: install a web cache


Scenario:

- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
 - avg data rate to browsers: 1.50 Mbps

Cost: web cache (cheap!)

Performance:


- LAN utilization: .?
 - access link utilization = ?
 - average end-end delay = ?
- How to compute link utilization, delay?*



Calculating access link utilization, end-end delay with cache:



suppose cache hit rate is 0.4:

- 40% requests served by cache, with low (msec) delay (same access network)
- 60% requests satisfied at origin
 - rate to browsers over access link
 $= 0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$
 - access link utilization $= 0.9/1.54 = .58$ means low (msec) queueing delay at access link



lower average end-end delay than with 154 Mbps link (and cheaper too!)

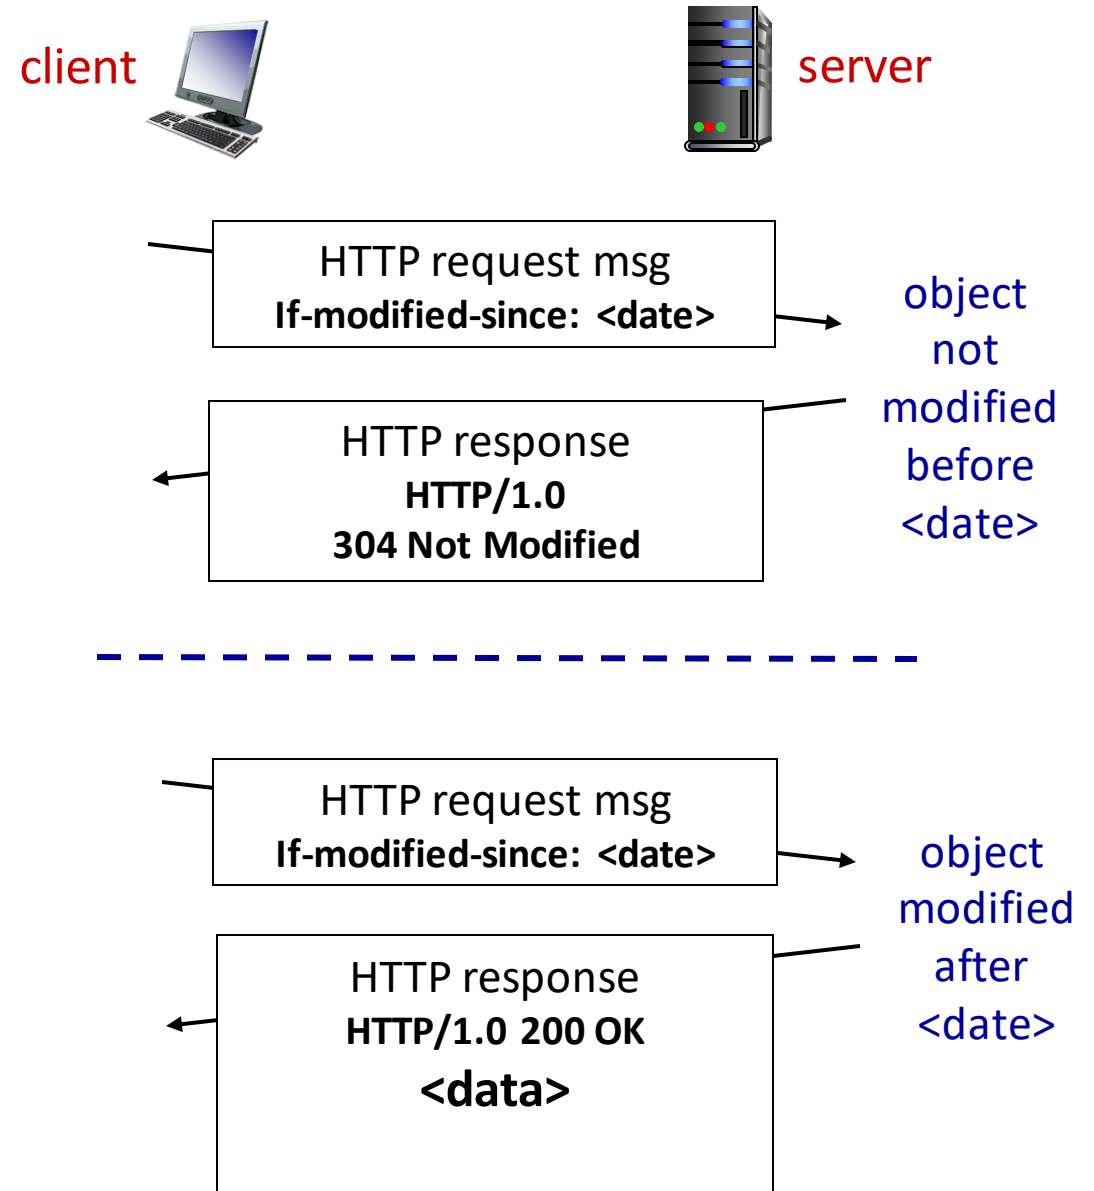
Web caches (cont.)



Conditional GET

Goal: don't send object if *client* has up-to-date cached version

- no object transmission delay (or use of network resources)
- *client*: specify date of cached copy in HTTP request
If-modified-since: <date>
- *server*: response contains no object if cached copy is up-to-date:
HTTP/1.0 304 Not Modified



HTTP/2

Key goal: decreased delay in multi-object HTTP requests

HTTP1.1: introduced multiple, pipelined GETs over single TCP connection

- server responds *in-order* (FCFS: first-come-first-served scheduling) to GET requests
- with FCFS, small object may have to wait for transmission (**head-of-line (HOL) blocking**) behind large object(s)
- loss recovery (retransmitting lost TCP segments) stalls object transmission

HTTP/2

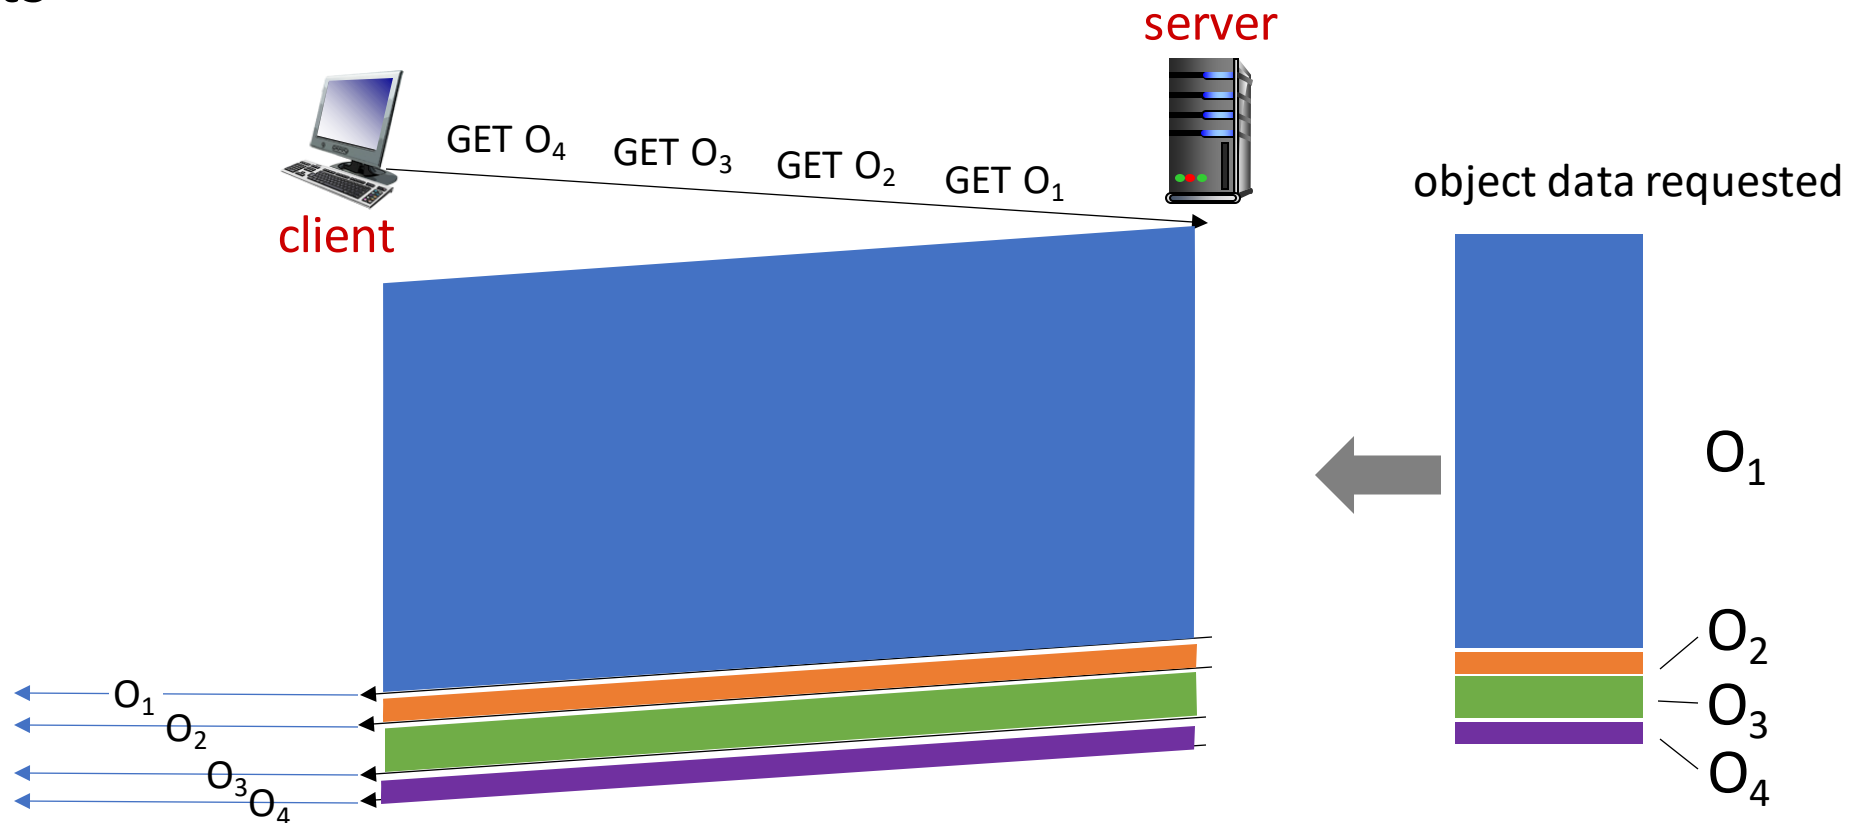
Key goal: decreased delay in multi-object HTTP requests

HTTP/2: [RFC 7540, 2015] increased flexibility at *server* in sending objects to client:

- methods, status codes, most header fields unchanged from HTTP 1.1
- transmission order of requested objects based on client-specified object priority (not necessarily FCFS)
- *push* unrequested objects to client
- divide objects into frames, schedule frames to mitigate HOL blocking

HTTP/2: mitigating HOL blocking

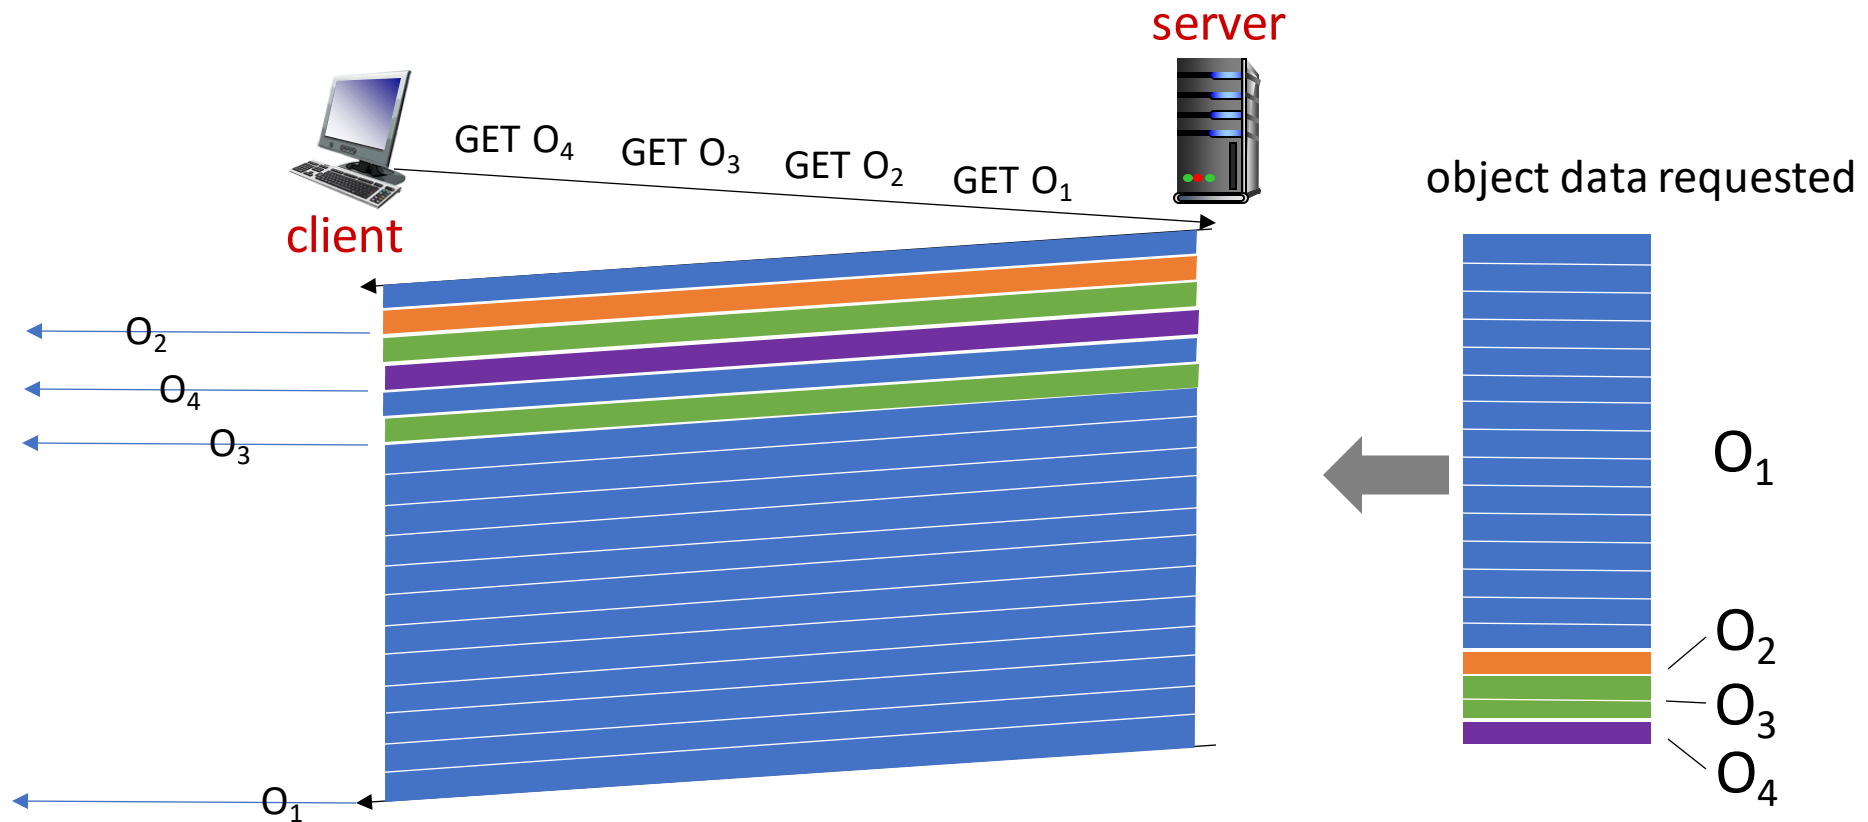
HTTP 1.1: client requests 1 large object (e.g., video file) and 3 smaller objects



objects delivered in order requested: O_2 , O_3 , O_4 wait behind O_1

HTTP/2: mitigating HOL blocking

HTTP/2: objects divided into frames, frame transmission interleaved



O_2 , O_3 , O_4 delivered quickly, O_1 slightly delayed

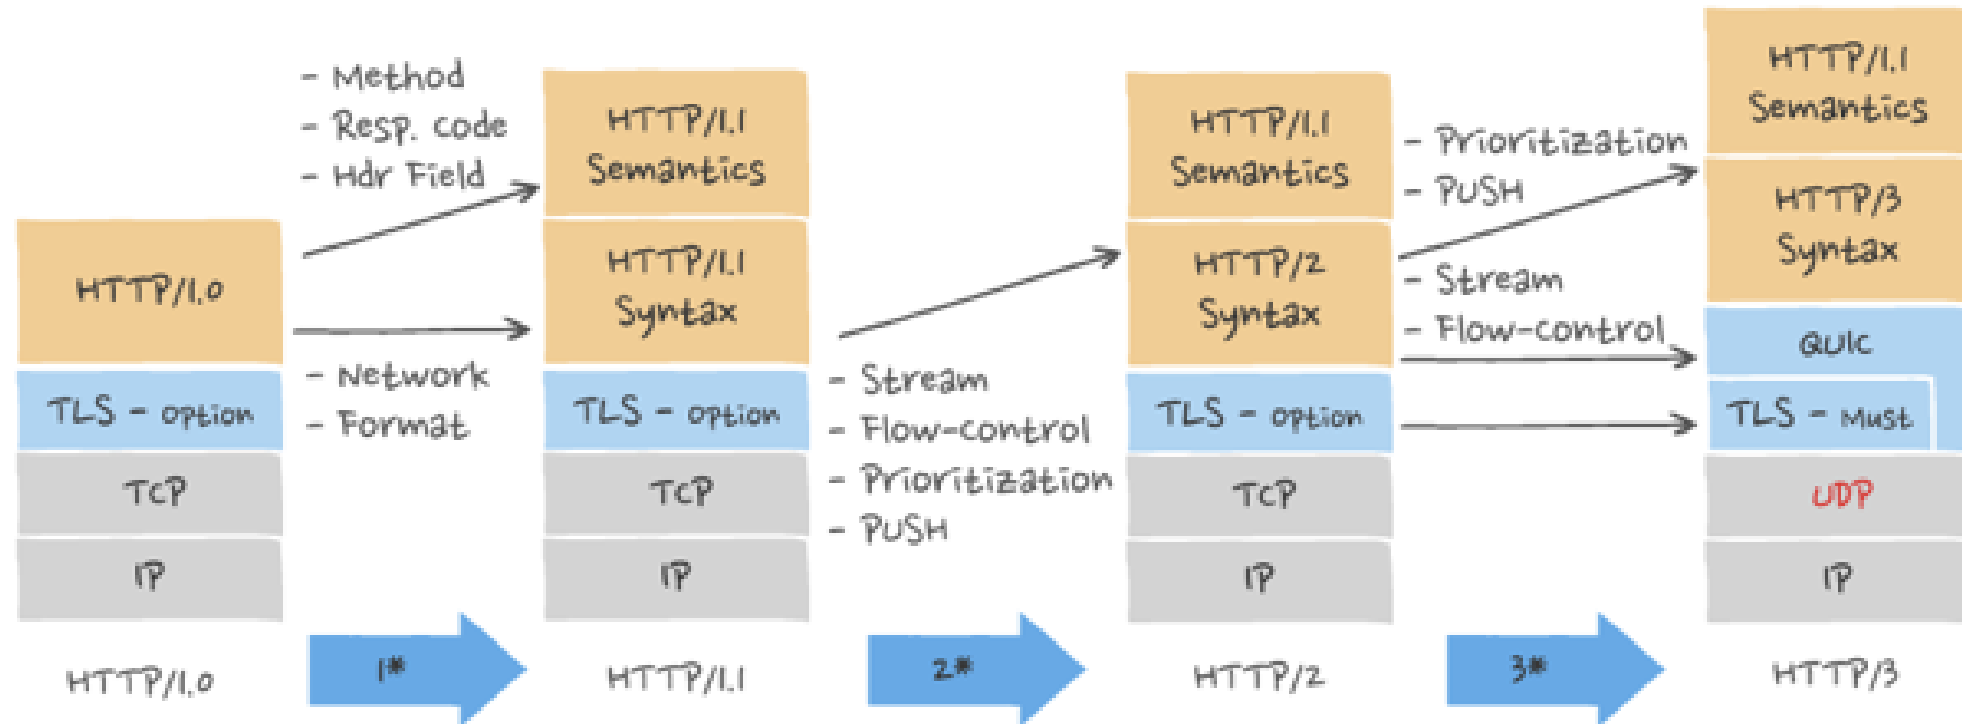
HTTP/2 to HTTP/3

HTTP/2 over single TCP connection means:

- recovery from packet loss still stalls all object transmissions
 - as in HTTP 1.1, browsers have incentive to open multiple parallel TCP connections to reduce stalling, increase overall throughput
- no security over vanilla TCP connection
- **HTTP/3**: adds security, per object error- and congestion-control (more pipelining) over UDP
 - more on HTTP/3 in transport layer

HTTP/1 vs HTTP/2 vs. HTTP/3

HTTP protocol stack transition and comparison



A H3/QUIC History

Launched by Google in 2012, support in Chrome in 2013:

- Google reported improvements in application performance
 - Reduction in YouTube video rebuffering and Google search latency
- >35% of Google's egress traffic (~7% of the internet) in 2017
- >75% of FB's traffic in 2020
- Check [here](#) for the latest figures

Facebook is bringing QUIC to billions

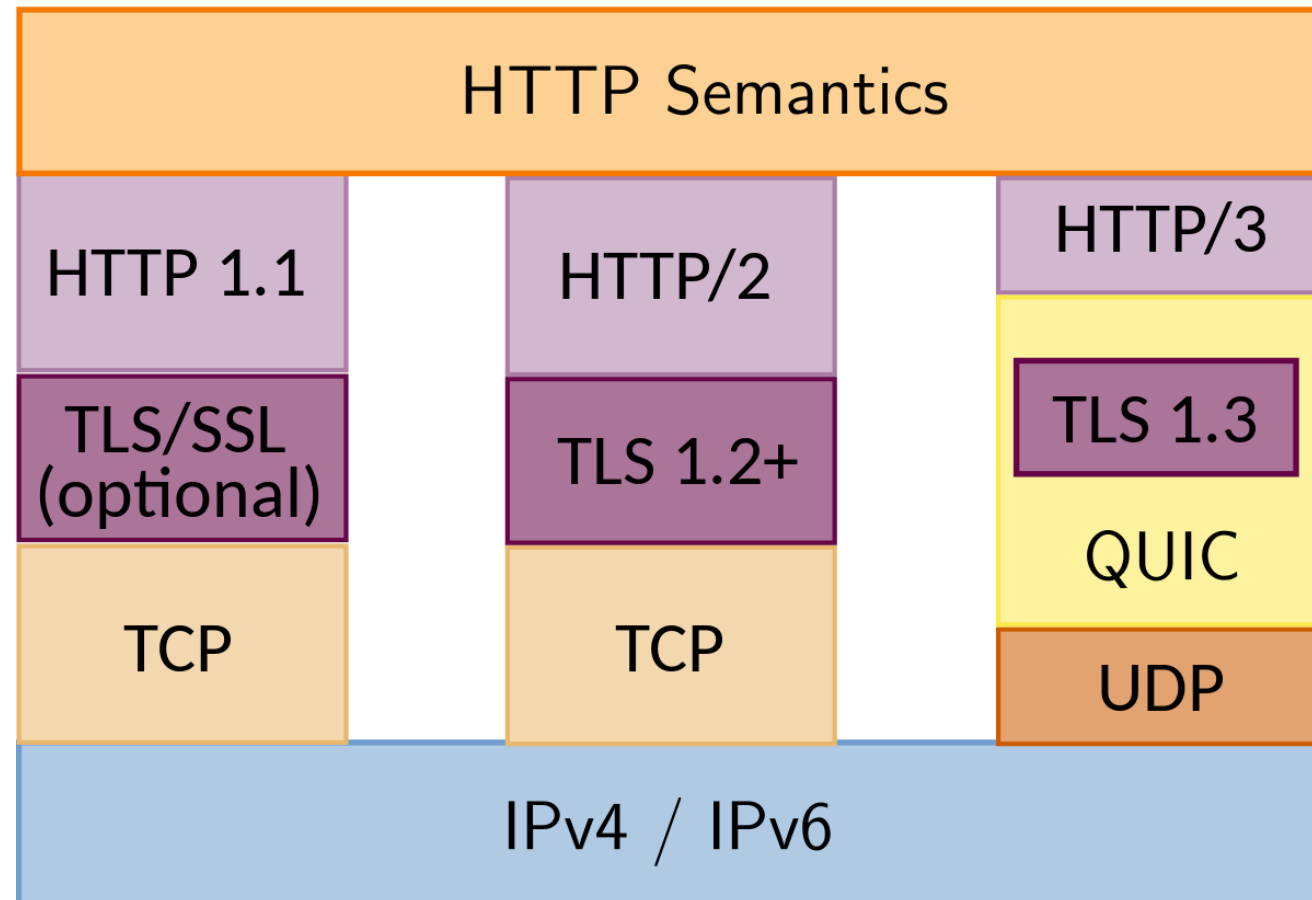
- QUIC for all content in the Facebook products
- QUIC had a transformative effect on video metrics in the Facebook app (e.g., stall rate was reduced by 20%)

QUIC Standardization: Six core documents are now with RFC editor

H3 (and QUIC) Support as of Today

Edge *	Firefox	Chrome	Safari	Opera	Safari on iOS *	Opera Mini *	Android Browser *	Opera Mobile *	Chrome for Android
12-18		4-78							
² 79-84	2-71	² 79-84		10-72					
³ 85-86	¹ 72-87	³ 85-86	3.1-13.1	³ 73	3.2-13.7				
87-91	88-90	87-91	⁴ 14	74-77	14.4		2.1-4.4.4	12-12.1	
92	91	92	⁴ 14.1	78	14.7	all	92	64	92
	92-93	93-95	⁴ 15-TP						

TCP/IP Protocol Stack (Until L3)



H3/QUIC Features

Main new features over TCP+TLS+H2:

- Connection establishment latency: 0-RTT (or 1-RTT)
- Customizable congestion control, improved retransmission machinery, forward error correction (FEC)
- Multiplexing w/o HoL blocking
- Connection migration: Moving between network interfaces without renegotiating the session

Reliable and prioritized delivery

- Decoupled retransmissions, congestion control and flow control
- Stream prioritization is managed by the sender

Encrypted delivery

- The network cannot identify the streams, so multiple connections are needed for QoS