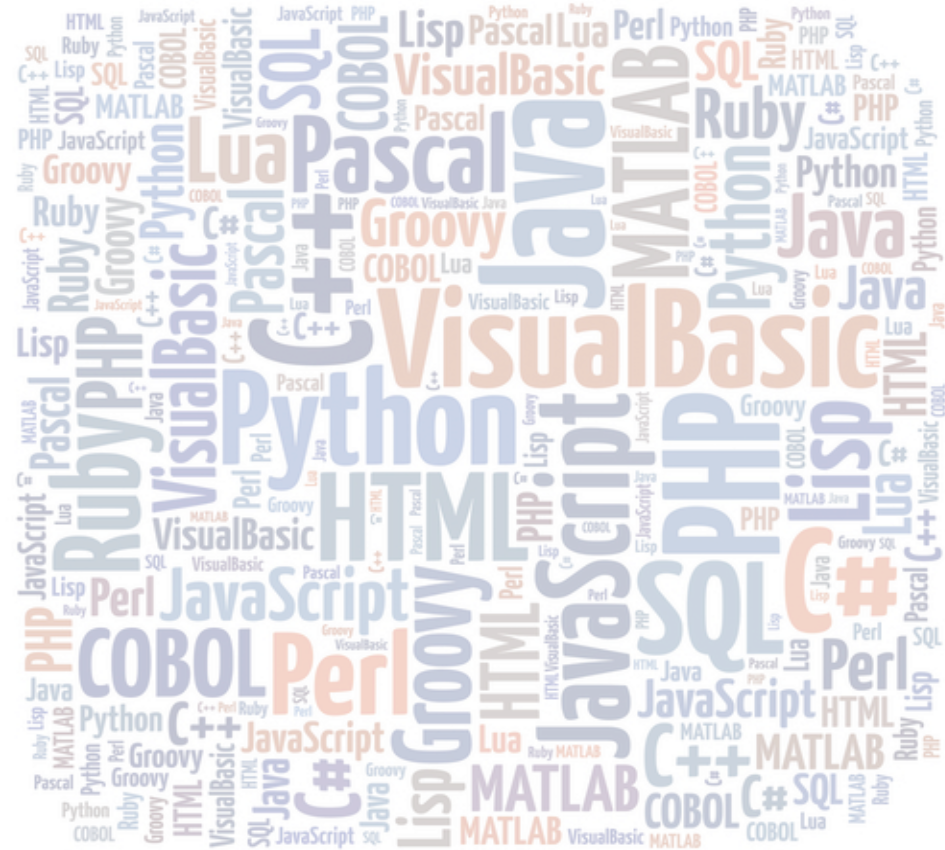


Composite Data Structures



Introduction

- Clojure data structures
 - Maps
 - Vectors
 - Lists
 - Sets
- Creation and access methods
- Immutability
 - How it's done in practice

Managing data

- In the last section, we looked at simple Clojure types
 - ints, floats, strings, and booleans
- But we can also work with powerful, composite data types.
- Clojure includes:
 - Maps
 - Vectors
 - Lists
 - Sets

Maps

- Clojure maps are the equivalent of Python dictionaries.
 - That is, they provide key/value lookup pairs.
- An empty map can be declared as { }
- To define a map with a collection of key/value pairs, we can do the following

```
; profession map
{
  "John"  "professor"
  "Sue"   "doctor"
  "Ahmed" "astronaut"
}
```

Maps...cont'd

- We can also do these things

```
; number map with int keys
{
  1 "one"
  2 "two"
  3 "three"
}
```

```
; function map
{
  "func1" (fn [] (println "func1"))
  "func2" (fn [] (println "func2"))
}
```

Keys in practice

- In practice, Clojure keys are often created using what are known as keywords.
- A keyword is a label prefixed with a ":"
- In a map definition, we might have:

```
; profession map
{
  :john "professor"
  :sue  "doctor"
  :ahmed "astronaut"
}
```

- Keywords are essentially strings, but provide faster internal lookups
 - Keywords are case sensitive
 - You will see them a lot

Map lookups

- Of course, we need to be able to find things in a map.
- For this we use the `get` function
 - Note that in many cases, you will want to provide a name for your maps in order to perform the lookup

```
(def jobs {  
  :john "professor"  
  :sue "doctor"  
  :ahmed "astronaut"  
})
```

```
(get jobs :sue) ; => "doctor"  
(get jobs :Sue) ; => nil
```

...and a little more

- Now let's put a few of these things together

```
; function lookup table
(def funcs {
  :func1 (fn [] (println "func1"))
  :func2 (fn [] (println "func2"))
  :func3 (fn [] (println "func3"))
})

; execute the third function
((get funcs :func3)) ; => func3
```

- What is the return value?
 - nil (from the println)

Odds and Ends

- Maps can be nested
 - The `get-in` function can access the inner map(s)

```
(def nest {
  :a "eh"
  :b {:dog "fido" :cat "fluffy"}
})

(get nest :a) ; => "eh"
(get nest :b)
  ; => {:dog "fido", :cat "fluffy"}

(get-in nest [:b :cat]) ; => "fluffy"
```

Odds and ends...cont'd

- You can provide default values for lookups
 - (get jobs :bill “not found”)
 - Returns “not found” if :bill is not in the map
- You can also use keywords like a function invocation (instead of get)
 - (:john jobs)
 - (:john jobs “not found”)
- There are other useful map functions, including, but not limited to
 - (keys map_name)
 - (vals map_name)
 - (contains? map_key)

One last thing

- We have said that Clojure uses whitespace to separate elements of the language
- While this is true, it is worth pointing out that commas are considered to be whitespace by Clojure
 - In other words, it simply ignores them
- However, you can sometimes use commas to make your code slightly more readable.
 - This is sometimes done with maps

```
{:john "prof" :sure "doctor"}
```

```
{:john "prof", :sure "doctor"}
```

Vectors

- A vector is essentially a 0-indexed sequence (array).
 - It is delimited by [] brackets
- Creating a vector is easy, using either of two techniques:

```
(def myVec [1 2 3])
```

```
(def myVec (vector 1 2 3))
```

- Again, commas could be used but this would not normally be done

Vectors...cont'd

- Again, we can use the `get` function to retrieve elements.
 - `nil` is returned if the index is out of bounds

```
(def things [  
  "dog"  
  2  
  (fn [] "do something")  
)  
  
(get things 0) ; => "dog"  
(get things 1) : => 2  
(get things 10) : => nil
```

- Note that we can put anything into a vector, including functions.

Lists

- As one would expect, Clojure provides a *list* data structure.
- Lists are delimited by the () parentheses
- If you want to define a list using the () notation, however, you have to prefix the list with a single quote character. Why?
 - If you don't do this, Clojure will interpret the list as an expression and try to execute it (which will probably fail)

```
(1 2 3) ; => error: 1 is not a function
```

```
`(1 2 3) ; => valid list
```

```
(+ 2 2) ; => function call
```

```
`(+ 2 2) ; => valid list
```

List content

- Note that, like vectors, it is also possible to define a list using the (list *list_content*) form.
- In addition, one can add any type of data to a list, including lists, vectors and maps.

```
(def boo (list 1 "two" [1 2 3]))
```

```
(def foo `(2 {:a "eh", :b "bee"} 3))
```

Lists...cont'd

- Lists have no indexing
- The `get` function does not work with lists (it always returns `nil`)
- Instead, use the `nth` function
 - An *out of bounds* error is generated for invalid indexes

```
(nth '(1 2 3) 1) ; => 2
```

```
(nth '(1 2 3) 10) ; error
```


Lists...cont'd

- How fast is `nth`?
- Recall that searches through a list take time proportional to the length of the list
 - On average, this is $O(n)$
- This isn't a problem in itself, but you should be aware that `nth` on a list is slower than `get` on a vector
 - Often, this dictates which data structure you will use in practice.
 - There isn't a lot of difference between the two so the choice depends on how you would like to access and update the data.

Sets

- Like Python, Clojure has a *set* data structure
- Its purpose is to maintain a collection of unique values.
- The syntax use the `#{ }` notation
 - The `#` is necessary to distinguish a set from a map
- It is also possible to use the `(hash-set values)` form.

```
#{1 2 "dog" "cat"}
```

```
(hash-set 1 2 "dog" :c)
```

Sets...cont'd

- Sets will automatically discard duplicates

```
# {1 2 "dog" 2} ; => # {1 2 "dog" }
```

- We can also create sets from existing lists and vectors, eliminating duplicates in the process

```
(set ` ( 1 1 2 2) ) ; => # {1 2}
```

```
(set ["dog" "cat" "dog"]) ; => # {"dog" "cat" }
```

Searching the set

- There are in fact three ways to search for an element in a set
 - get
 - Using a keyword as a function
 - contains?

```
(def foo #{ 1 2 3 :a} )

(get foo 3)           ; => 3
(get foo 4)           ; => nil

(:a foo)              ; => :a, returns itself

(contains? foo 2)     ; => true
```

Immutability

- Clojure data structures are immutable.
 - That is, one does not actually change the values within a structure.
- Specifically, when we make “changes”, we are actually creating a new/second data structure.
 - Same as the first, except with the modification.
- You can see this simply by using the REPL interface to modify an existing structure, and then checking the structure again.

Adding to a sequence

- To show this, note that Clojure provides a function called `conj` (i.e., conjoin) to add a value to lists and vectors
- In practice, `conj` adds a value to the *beginning* of a list and the *end* of a vector. Why?

```
(def foo1 '(1 2 3))  
(conj foo1 0)      ; => (0 1 2 3)  
foo1                ; => (1 2 3), no change  
  
(def foo2 (conj foo1 0))  
foo1                ; => (1 2 3)  
foo2                ; => (0 1 2 3)
```

Efficiency

- What if I need to make thousands of changes to a large list
- I would now have thousands of versions of the large list
 - YUCK!!!
- This could be catastrophic in terms of:
 - memory consumption
 - processing performance (all of that copying)
- Is Clojure (and other functional languages) really that bad?

Efficiency...cont'd


- It is true that Clojure does make copies of the data structures.
- But it does this in a very clever way.
- We will use the example of a vector, but the principle can be applied to other data structures
- Before looking at the model, we will note that the `assoc` function can be used to update a specific vector position
 - Again, the “update” does not actually change the original vector

```
(def foo [1 2 3])  
foo ; => [1 2 3]
```

```
(assoc foo 1 "dog") ; => [1 "dog" 3]  
foo ; => [1 2 3]
```

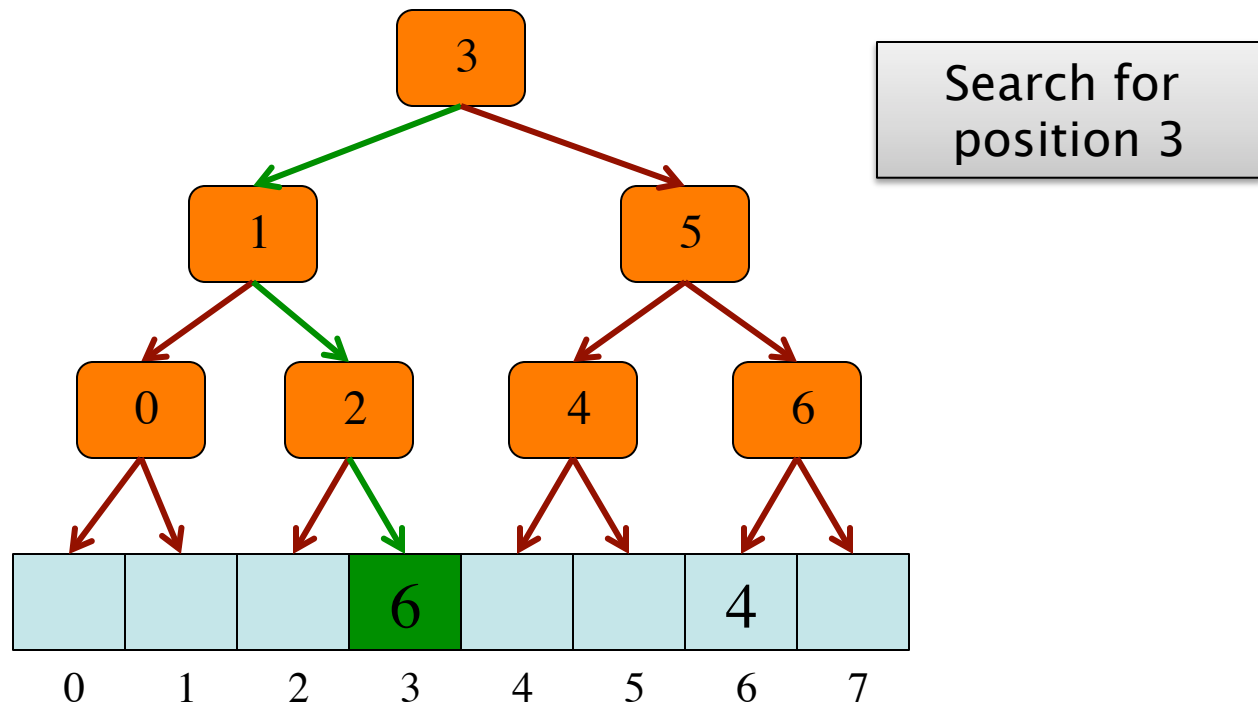

Modifications

- Let's start by looking at a basic vector implementation
- Below, we have a simple array
- In this case, 2 of the 8 cells have a non-nil value
- Let's say that we want to change the 6 to a 9.
 - Clearly, we could modify the 3rd cell in place but this would destroy the immutability property
 - We could also create a new array with the updated value, but this would be very inefficient



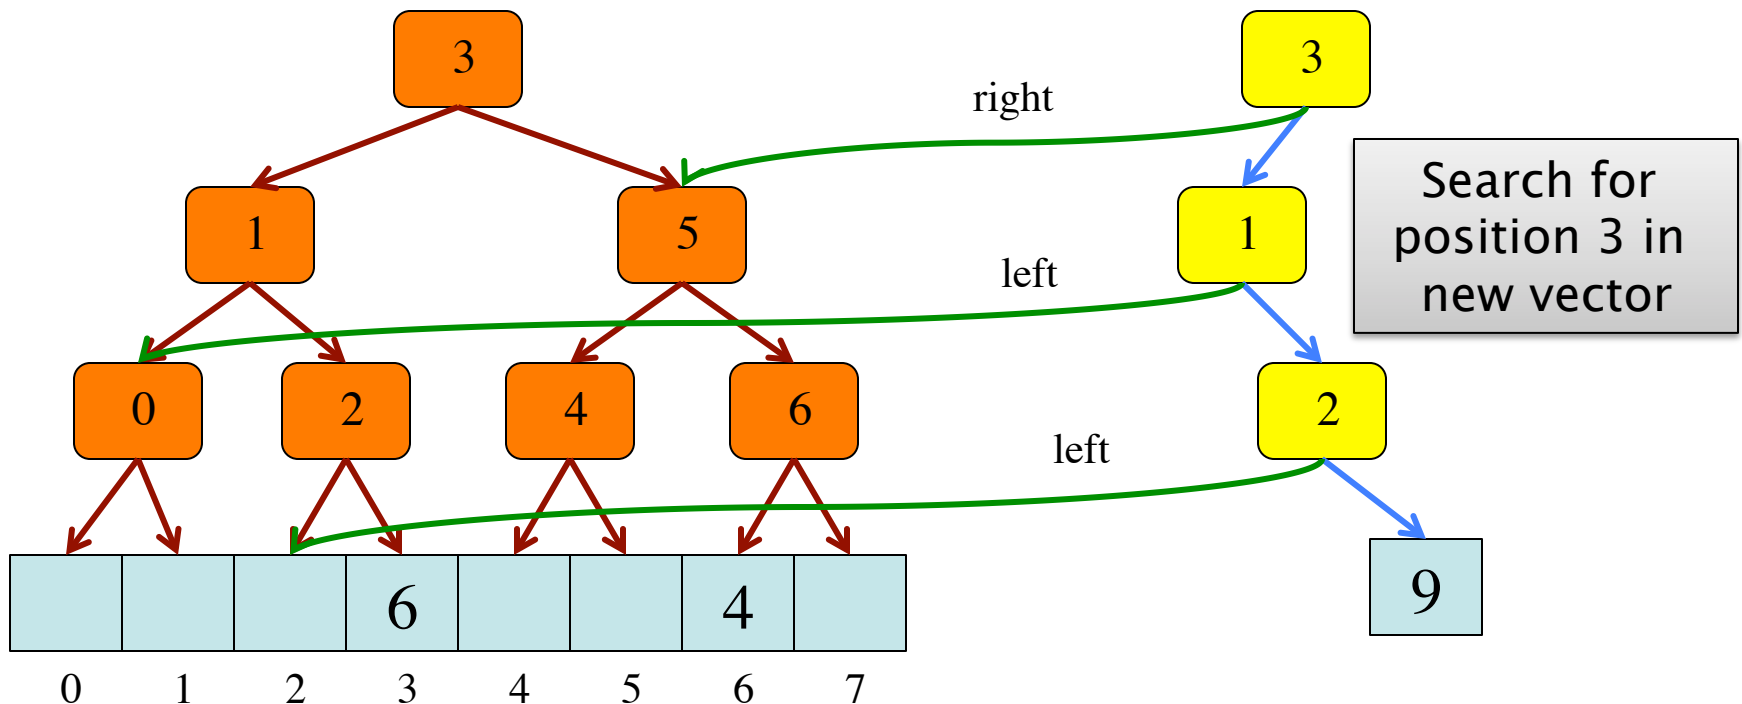
Modification...cont'd

- Instead, let's augment the basic data structure itself by including an indexing structure
- The index will be traversed by going left if the required position is \leq the node value, and right otherwise



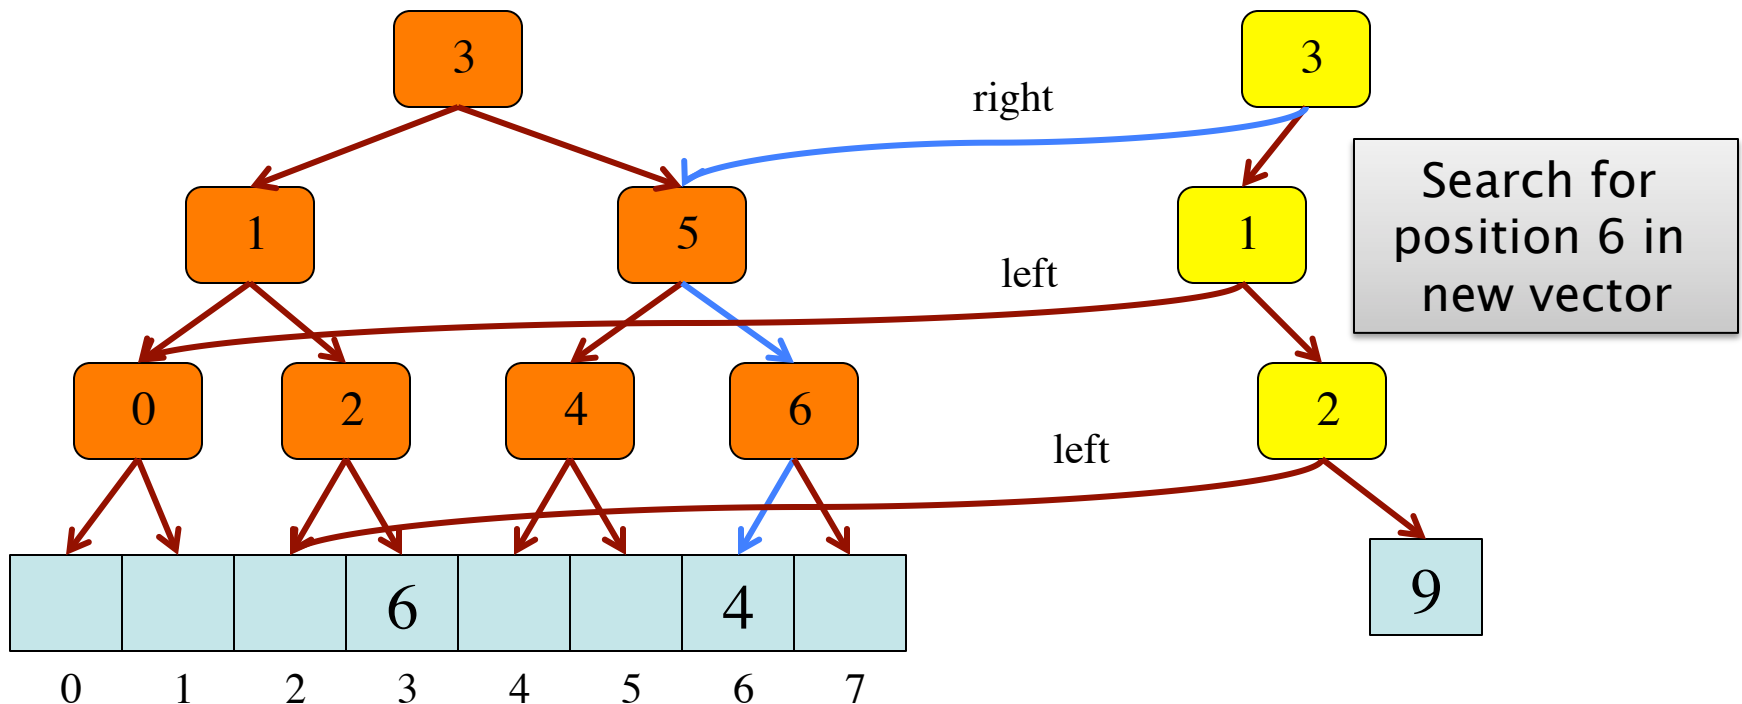
Modification...cont'd

- Now let's update 6 to 9
- To do this, we will create a new data structure that simply consists of a new path for the updated element
- We will then connect the remaining paths to the original index
- To find the updated element, we will use the new index (see path in blue)



Searching the new vector

- Now let's search for position 6 in the new vector.
- In this case, we navigate back to the first data structure and follow the index to position 6 (see path in blue).



Modification...cont'd

- Now, we have two distinct data structures, with the second allowing access to the updated elements.
- However, for all other elements, the unmodified elements are shared.
- In effect, this creates a core *read-only* data structure that is modified as required.
- If we now modify the second data structure, a third data structure can point back to #2 or to #1.

Modifications...cont'd

- What does this cost?
- Without going into a great deal of theory, the indexes can be “balanced”, so that they are essentially shaped like a triangle.
- Such trees have logarithmic (\log_2) access time.
 - For example, an index with 1000 nodes has a height of just 10, and a tree with 1M nodes has a height of just 20.
- Since logarithms grow VERY slowly, access times are very good, though not of course quite as fast as the $O(1)$ times of the original vector.