

COMP 6721 Applied Artificial Intelligence (Summer 2023)

Assignment #5: Decision Trees & Evaluation Metrics

Question 1 Given the training instances below, use information theory to find whether ‘Outlook’ or ‘Windy’ is the best feature to decide when to play a game of golf.

Outlook	Temperature	Humidity	Windy	Play / Don't Play
sunny	85	85	false	Don't Play
sunny	80	90	true	Don't Play
overcast	83	78	false	Play
rain	70	96	false	Play
rain	68	80	false	Play
rain	65	70	true	Don't Play
overcast	64	65	true	Play
sunny	72	95	false	Don't Play
sunny	69	70	false	Play
rain	75	80	false	Play
sunny	75	70	true	Play
overcast	72	90	true	Play
overcast	81	75	false	Play
rain	71	80	true	Don't Play

Question 2 It's time to leave the calculations to your computer: Write a Python program that uses `scikit-learn`'s *Decision Tree Classifier*.¹

```
import numpy as np
from sklearn import tree
from sklearn import preprocessing
```

Here is the training data from the first question:

```
dataset = np.array([
    ['sunny', 85, 85, 0, 'Don\'t Play'],
    ['sunny', 80, 90, 1, 'Don\'t Play'],
    ['overcast', 83, 78, 0, 'Play'],
    ['rain', 70, 96, 0, 'Play'],
    ['rain', 68, 80, 0, 'Play'],
    ['rain', 65, 70, 1, 'Don\'t Play'],
    ['overcast', 64, 65, 1, 'Play'],
    ['sunny', 72, 95, 0, 'Don\'t Play'],
    ['sunny', 69, 70, 0, 'Play'],
    ['rain', 75, 80, 0, 'Play'],
    ['sunny', 75, 70, 1, 'Play'],
    ['overcast', 72, 90, 1, 'Play'],
    ['overcast', 81, 75, 0, 'Play'],
    ['rain', 71, 80, 1, 'Don\'t Play'],
])
```

Note that we changed `True` and `False` into `1` and `0`.

For our feature vectors, we need the first four columns:

```
X = dataset[:, 0:4]
```

and for the training labels, we use the last column from the dataset:

```
y = dataset[:, 4]
```

However, you will not be able to use the data as-is: All features must be numerical for training the classifier, so you have to transform the strings into numbers. Fortunately, `scikit-learn` has a preprocessing class for *label encoding* that we can use:²

```
le = preprocessing.LabelEncoder()
X[:, 0] = le.fit_transform(X[:, 0])
```

(Note: you will need to transform `y` as well.)

¹See <https://scikit-learn.org/stable/modules/tree.html#classification>

²See https://scikit-learn.org/stable/modules/preprocessing_targets.html#preprocessing-targets

Now you can create a classifier object:

```
dtc = tree.DecisionTreeClassifier(criterion="entropy")
```

Note that we are using the *entropy* option for building the tree, which is the method we've studied in class. Train the classifier to build the tree:

```
dtc.fit(X, y)
```

Now you can predict a new value using `dtc.predict(test)`, just like you did for the Naïve Bayes classifier last week. Note: if you want the string output that you transformed above, you can use the `inverse_transform(predict)` function of a label encoder to change the `predicted` result back into a string.

You can also print the tree:

```
tree.plot_tree(dtc)
```

but this can be a bit hard to read; to get a prettier version you can use the *Graphviz* visualization software,³ which you can call from Python like this:

```
# print a nicer tree using graphviz
import graphviz
dot_data = tree.export_graphviz(dtc, out_file=None,
    feature_names=['Outlook', 'Temperature', 'Humidity', 'Windy'],
    class_names=['Don\'t Play', 'Play'],
    filled=True, rounded=True)
graph = graphviz.Source(dot_data)
graph.render("mytree")
```

The result will be stored in a file `mytree.pdf` and should look like Figure 1.

³See <https://www.graphviz.org/>

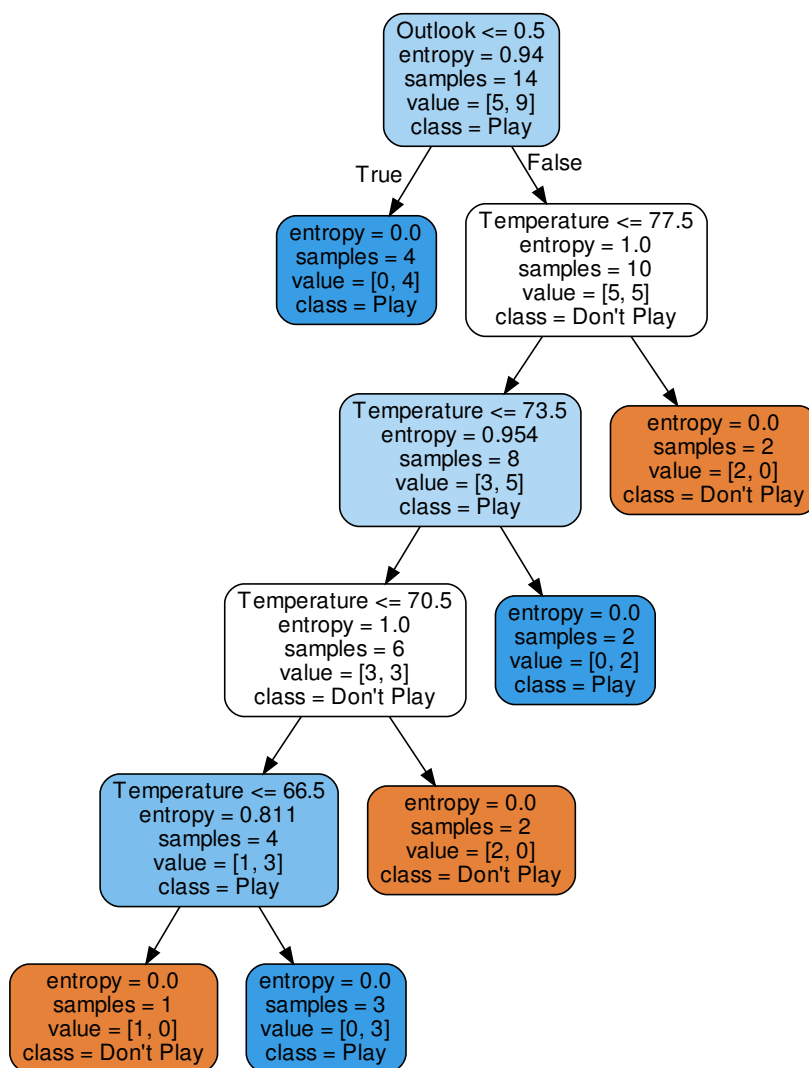


Figure 1: Generated Decision Tree using scikit-learn: Note that the string values for *Outlook* have been changed into numerical values ('overcast'= 0, 'rain'= 1, 'sunny' = 2)

Question 3 For the tree of the last question, report the following parameters or metrics:

- (a) Number of leaves
- (b) Height (depth) of the tree
- (c) External Path Length
- (d) Weighted External Path Length
- (e) Accuracy, recall, and precision
- (f) Now assume that we do pruning considering the hyper parameter of "Minimum Leaf Size" =2. What will be the accuracy, recall, and precision of this tree (training)?

Question 4 Let's start working with some “real” data: `scikit-learn` comes with a number of example datasets, including the *Iris* flower dataset. If you do not know this dataset, start by reading https://en.wikipedia.org/wiki/Iris_flower_dataset.⁴

- (a) Change your program from the previous question to work with this new dataset:

```
from sklearn.datasets import load_iris
X, y = load_iris(return_X_y=True)
```

Train a decision tree and print it out like before.

- (b) Now you have to *evaluate* the performance of your classifier. Use `scikit-learn`'s `train_test_split` helper function⁵ to split the Iris dataset into a *training* and *testing* subset, as discussed in the lecture. Now run an *evaluation* to compute the *Precision*, *Recall*, *F₁-measure*, and *Accuracy* of your classifier using the evaluation tools in `scikit-learn`.⁶
- (c) Finally, compute and print out the *confusion matrix*.⁷



Figure 2: An *Iris versicolor*⁸

⁴Also see https://scikit-learn.org/stable/auto_examples/datasets/plot_iris_dataset.html for the `scikit-learn` documentation of this dataset.

⁵See https://scikit-learn.org/stable/modules/cross_validation.html

⁶See https://scikit-learn.org/stable/modules/model_evaluation.html

⁷See https://scikit-learn.org/stable/modules/generated/sklearn.metrics.confusion_matrix.html

⁸Photo taken by Danielle Langlois in July 2005 at the Forillon National Park of Canada, Quebec, Canada. Licensed under the Creative Commons Attribution-Share Alike 3.0 Unported license. https://commons.wikimedia.org/wiki/File:Iris_versicolor_3.jpg