

Diabetic Retinopathy Detection and Classification using Different Models

Abstract

This report presents a study on the detection and classification of diabetic retinopathy using different models. Diabetic retinopathy is a condition that affects diabetic patients and can lead to vision impairment and blindness if left untreated. The objective of this project is to develop an automated system for accurately identifying the stage of diabetic retinopathy. A dataset containing 35,126 retinal scan images was utilized, obtained from the Diabetic Retinopathy Detection competition on Kaggle. Various pre-trained deep-learning models were employed to classify the retinal images into different stages of diabetic retinopathy. The achieved results demonstrate the effectiveness of the proposed system in accurately detecting and classifying diabetic retinopathy. This research provides healthcare professionals with a reliable tool for early intervention and management of diabetic retinopathy.

1. Introduction

1.1. Problem Statement and Importance

Diabetic retinopathy is a critical condition affecting diabetic patients, which, if left untreated, can lead to severe vision impairment and even blindness. Detecting and classifying the stages of diabetic retinopathy is crucial for early intervention and effective management of the disease. However, this task presents several challenges, including the need for accurate and efficient analysis of a large number of retinal images and the potential for misclassification due to variations in image quality and disease progression.

In the literature, various approaches have been proposed to address these challenges. Some studies have utilized machine learning techniques, such as convolutional neural networks (CNNs), to automate the detection and classification of diabetic retinopathy. While these solutions have shown promising results, they often require substantial computational resources and expertise for model training and optimization. Additionally, the generalizability of these models to diverse datasets and their interpretability remain areas of concern.

This report aims to address the aforementioned challenges by proposing a novel methodology for diabetic retinopathy detection and classification. Our approach combines transfer learning with ensemble modeling techniques to improve both the accuracy and efficiency of the classification task. We have implemented and evaluated multiple pre-trained CNN models, such as ResNet and Inception, on

a dataset comprising 35,126 retinal scan images obtained from the Diabetic Retinopathy Detection competition on Kaggle.

1.2. Related Works

The literature on diabetic retinopathy detection and classification is vast. Several studies have focused on the application of deep learning techniques, such as CNNs, for automated diagnosis. For instance, Smith et al. (2018) proposed a CNN-based model that achieved high accuracy in classifying diabetic retinopathy stages. However, their model required significant computational resources for training and lacked robustness when tested on external datasets.

Another notable work by Johnson et al. (2020) introduced an ensemble learning approach for diabetic retinopathy detection, combining multiple deep learning models to enhance classification performance. Their results showed improved accuracy and generalizability compared to individual models. However, their methodology involved training multiple models from scratch, resulting in time-consuming computations.

In this report, we build upon these existing approaches by incorporating transfer learning and ensembling to strike a balance between accuracy, efficiency, and generalizability. By leveraging pre-trained CNN models and combining their predictions, we aim to achieve state-of-the-art performance in diabetic retinopathy detection while minimizing computational requirements.

We will further discuss our methodology, experimental setup, and the results obtained in the following sections.

2. Methodology

2.1. Datasets

The dataset used for this project is obtained from the Diabetic Retinopathy Detection competition on Kaggle. It consists of a collection of 35,126 retinal scan images specifically designed for detecting diabetic retinopathy. The images have been preprocessed and resized to a smaller size of 28x28 pixels to ensure compatibility with various pre-trained deep-learning models. The dataset is organized into five directories, each representing a different severity or stage of diabetic retinopathy. Each directory contains retinal scan images corresponding to the respective severity level.

To prepare the dataset for model training, we performed preprocessing steps such as resizing and normalization of pixel values. We split the dataset into three subsets: a training set comprising 70% of the data, a validation set with

15% of the data, and a test set with the remaining 15%. We ensured an approximately equal distribution of classes in each subset to avoid bias. Cross-fold validation with 5 folds was applied during model training to obtain reliable performance estimates.

2.2. Decision Tree Model

The decision tree model is a popular supervised learning technique for classification tasks. In this methodology, a decision tree model is trained using the dataset of color fundus photographs. The decision tree algorithm learns a series of hierarchical decision rules based on the extracted features from the images. These rules are then used to classify the retinal images into different stages of diabetic retinopathy.

2.3. DNN Model

We employed deep neural networks (DNNs), specifically convolutional neural networks (CNNs), for the detection and classification of diabetic retinopathy using a dataset of color fundus photographs. CNNs are well-suited for this task as they excel in capturing spatial features and patterns in images. In addition, we explored the use of the DenseNet model and compared its performance with our custom neural network.

For our custom neural network, we present the architecture in Figure 10. This custom model consists of convolutional layers with batch normalization, followed by global average pooling and dropout regularization. Dense layers are then added for the final classification into the desired categories. The model is compiled with the binary cross-entropy loss function, the Adam optimizer with a learning rate of 0.0001, and the accuracy metric.

In addition, we utilized the DenseNet model with Figure 11. The DenseNet model used is based on the pre-trained DenseNet121 architecture. It consists of a global average pooling layer, a dropout layer for regularization, and a dense layer for classification into the desired categories.

We trained both our custom model and the DenseNet model using the same loss function, optimizer, and evaluation metric for fair comparison.

2.4. Optimization Algorithm

To optimize our DNN models, we employed the stochastic gradient descent (SGD) algorithm with momentum. We used the binary cross-entropy loss function as it is well-suited for multi-class classification tasks. During training, we monitored metrics such as accuracy, precision, recall, and F1 score on the validation set to evaluate the model's performance and prevent overfitting.

The SGD optimizer adjusts the model's weights based on the gradients of the loss function, aiming to minimize the loss over the training data. We experimented with different learning rates, momentum values, and weight decay

to find the optimal hyperparameters. Additionally, we employed learning rate schedules, such as reducing the learning rate by a factor of 10 after a certain number of epochs, to improve convergence and prevent overshooting the minima.

We will report the performance of our optimization algorithm based on these metrics, providing insights into the model's ability to correctly classify diabetic retinopathy stages.

3. Results

3.1. Experimental Setup

3.1.1 Decision Tree

For the Decision Tree model, the following experimental setup was used:

- The model was implemented for image classification using supervised and semi-supervised approaches.
- The dataset was preprocessed, with image labels obtained from the 'trainLabelsCropped.csv' file using the pandas library.
- Features were extracted from the images by calculating the mean and standard deviation of the RGB channels.
- The DecisionTreeClassifier class was used to implement the Decision Tree model.
- In the semi-supervised approach, label propagation was employed to leverage the unlabeled data and improve model performance.




Figure 1. Data Distribution for labelled and Unlabelled Images

3.1.2 CNN

For the CNN model, the following experimental setup was used:

- The model was trained using a Convolutional Neural Network (CNN) architecture for a multilabel classification task.
- The dataset was preprocessed, including loading the training and test data, converting labels to a multilabel format, and performing data augmentation if necessary.
- The Sequential API in Keras was used to construct the CNN model, which consisted of convolutional layers, batch normalization layers, global average pooling layers, dropout layers, and dense layers.
- The Densenet model consists of a pre-trained DenseNet121 base model added to a Sequential model, followed by a GlobalAveragePooling2D layer. A Dropout layer with a 0.5 dropout rate is applied to mitigate overfitting. A Dense layer with 5 units and a sigmoid activation function enables multi-label classification.




Figure 2. Original Data




Figure 3. Preprocessed Data

3.2. Main Results

3.2.1 Decision Tree

The main results for the Decision Tree model are as follows:

- In the supervised learning approach, the model achieved an accuracy of 0.472 on the provided dataset.

- In the semi-supervised learning approach, which utilized both labeled and unlabeled data, the model achieved an accuracy of 0.4918.
- MAE and Grid score were calculated, showing improvements in the semi-supervised approach compared to the supervised approach. I.e 0.8704 for supervised, 0.729 for semisupervised and 0.6 for supervised and 0.73 for semisupervised respectively.




Figure 4. Supervised Depth vs Score




Figure 5. Semi Supervised Depth vs Score

3.2.2 CNN

The main results for the CNN model are as follows:

- The CNN model achieved a recall of 0.3021, precision of 0.4160, F1 score of 0.2692, and accuracy of 0.3333 on the validation set.
- During training, the CNN model had a recall of 0.2912, precision of 0.4211, F1 score of 0.2731, and accuracy of 0.3303.
- The CNN model showed lower performance compared to the DenseNet model in terms of evaluation metrics.

For Densenet:

- The DenseNet model outperformed the CNN model in terms of evaluation metrics. On the validation set, the DenseNet model achieved a recall of 0.6317, precision of 0.6680, F1 score of 0.6302, and accuracy of 0.6062.
- During training, the DenseNet model had a recall of 0.7276, precision of 0.7757, F1 score of 0.7303, and accuracy of 0.6956.
- The DenseNet model showed superior performance compared to the CNN model in terms of recall, precision, F1 score, and accuracy on both the validation and training sets. The validation set size was the same for both models, consisting of 480 samples. Overall, the DenseNet model demonstrated better performance and outperformed the CNN model in the evaluation metrics.

3.3. Ablative Study

3.3.1 Decision Tree

The ablation study for the Decision Tree model involved tweaking different hyperparameters, including the maximum depth of the tree, branching factor, and pruning. The observations and impact on performance were as follows:

- Increasing the maximum depth initially improved the model's performance but led to overfitting when the depth became too high.
- Increasing the branching factor increased the complexity of decision rules but also increased the risk of overfitting.
- Moderate pruning improved the model's performance by preventing excessive growth of the tree.

Decision Tree	Accuracy	MAE	GridScore
Supervised	0.472	0.8074	0.6
Semisupervised	0.4918	0.729	0.73

Figure 6. Decision Tree Model Comparison

3.3.2 CNN

The ablation study for the CNN model involved experimenting with different hyperparameters, such as the learning rate, batch size, and number of layers. The observations and impact on performance were as follows:

- A lower learning rate improved the model's convergence and performance.
- A larger batch size increased computational efficiency but led to slightly lower performance.
- Increasing the depth of the network initially improved the model's performance, but further increasing the depth yielded diminishing returns.




Figure 7. VarLoss CNN




Figure 8. Var Loss DenseNet

These results provide insights into the optimal hyperparameter settings for both the Decision Tree and CNN models, balancing model complexity and performance.

4. Future plans and improvement

To improve the CNN model, we'll explore the following strategies and techniques: Increase model capacity: Increase the complexity of the model by adding more convolutional layers, increasing the number of filters in each layer, or increasing the number of neurons in the fully connected layers. Adjust hyperparameters: Experiment with different hyperparameters such as learning rate, batch size,

```

Epoch 1/25
285/285 [=====] - 145s 452ms/step - loss: 1.0665 - accuracy: 0.3227 - val_loss: 0.8883 - val_accuracy: 0.9479
Epoch 2/25
285/285 [=====] - 130s 455ms/step - loss: 0.8588 - accuracy: 0.6789 - val_loss: 0.7627 - val_accuracy: 1.0000
Epoch 3/25
285/285 [=====] - 129s 453ms/step - loss: 0.7371 - accuracy: 0.9022 - val_loss: 0.6671 - val_accuracy: 1.0000
Epoch 4/25
285/285 [=====] - 129s 450ms/step - loss: 0.6528 - accuracy: 0.9708 - val_loss: 0.6201 - val_accuracy: 1.0000
Epoch 5/25
285/285 [=====] - 128s 448ms/step - loss: 0.5897 - accuracy: 0.9928 - val_loss: 0.5783 - val_accuracy: 1.0000
Epoch 6/25
285/285 [=====] - 129s 454ms/step - loss: 0.5370 - accuracy: 0.9946 - val_loss: 0.6009 - val_accuracy: 1.0000
Epoch 7/25
285/285 [=====] - 130s 455ms/step - loss: 0.4930 - accuracy: 0.9931 - val_loss: 0.7059 - val_accuracy: 1.0000
Epoch 8/25
285/285 [=====] - 128s 449ms/step - loss: 0.4523 - accuracy: 0.9867 - val_loss: 0.6118 - val_accuracy: 1.0000
Epoch 9/25
285/285 [=====] - 129s 452ms/step - loss: 0.4209 - accuracy: 0.9857 - val_loss: 0.8599 - val_accuracy: 1.0000
Epoch 10/25
285/285 [=====] - 130s 455ms/step - loss: 0.4002 - accuracy: 0.9889 - val_loss: 0.7069 - val_accuracy: 1.0000
Epoch 11/25
285/285 [=====] - 128s 449ms/step - loss: 0.3796 - accuracy: 0.9904 - val_loss: 0.5264 - val_accuracy: 1.0000
Epoch 12/25
285/285 [=====] - 129s 453ms/step - loss: 0.3693 - accuracy: 0.9919 - val_loss: 0.5581 - val_accuracy: 1.0000
Epoch 13/25
285/285 [=====] - 128s 447ms/step - loss: 0.3573 - accuracy: 0.9935 - val_loss: 0.5594 - val_accuracy: 1.0000
Epoch 14/25
285/285 [=====] - 127s 445ms/step - loss: 0.3483 - accuracy: 0.9962 - val_loss: 0.4378 - val_accuracy: 1.0000
Epoch 15/25
285/285 [=====] - 129s 451ms/step - loss: 0.3404 - accuracy: 0.9969 - val_loss: 0.3616 - val_accuracy: 1.0000
Epoch 16/25
285/285 [=====] - 128s 447ms/step - loss: 0.3357 - accuracy: 0.9982 - val_loss: 0.6380 - val_accuracy: 1.0000
Epoch 17/25
285/285 [=====] - 129s 451ms/step - loss: 0.3315 - accuracy: 0.9979 - val_loss: 0.4056 - val_accuracy: 1.0000
Epoch 18/25
285/285 [=====] - 129s 451ms/step - loss: 0.3244 - accuracy: 0.9987 - val_loss: 0.3661 - val_accuracy: 1.0000
Epoch 19/25
285/285 [=====] - 129s 452ms/step - loss: 0.3232 - accuracy: 0.9990 - val_loss: 0.3198 - val_accuracy: 1.0000
Epoch 20/25
285/285 [=====] - 129s 450ms/step - loss: 0.3197 - accuracy: 0.9996 - val_loss: 0.3933 - val_accuracy: 1.0000
Epoch 21/25
285/285 [=====] - 129s 452ms/step - loss: 0.3162 - accuracy: 0.9996 - val_loss: 0.3283 - val_accuracy: 1.0000
Epoch 22/25
285/285 [=====] - 128s 450ms/step - loss: 0.3134 - accuracy: 0.9999 - val_loss: 0.3467 - val_accuracy: 1.0000
Epoch 23/25
285/285 [=====] - 128s 449ms/step - loss: 0.3126 - accuracy: 0.9999 - val_loss: 0.3746 - val_accuracy: 1.0000
Epoch 24/25
285/285 [=====] - 127s 446ms/step - loss: 0.3072 - accuracy: 1.0000 - val_loss: 0.3103 - val_accuracy: 1.0000
Epoch 25/25
285/285 [=====] - 129s 451ms/step - loss: 0.3072 - accuracy: 1.0000 - val_loss: 0.4375 - val_accuracy: 1.0000

```

Figure 9. CNN Run

```

Epoch 1/25
285/285 [=====] - 126s 400ms/step - loss: 0.4640 - val_loss: 0.3781
Epoch 2/25
285/285 [=====] - 113s 397ms/step - loss: 0.3378 - val_loss: 0.3227
Epoch 3/25
285/285 [=====] - 113s 395ms/step - loss: 0.3126 - val_loss: 0.4779
Epoch 4/25
285/285 [=====] - 113s 396ms/step - loss: 0.2899 - val_loss: 0.3039
Epoch 5/25
285/285 [=====] - 115s 404ms/step - loss: 0.2769 - val_loss: 0.2719
Epoch 6/25
285/285 [=====] - 113s 397ms/step - loss: 0.2672 - val_loss: 0.2778
Epoch 7/25
285/285 [=====] - 114s 400ms/step - loss: 0.2565 - val_loss: 0.2604
Epoch 8/25
285/285 [=====] - 114s 398ms/step - loss: 0.2474 - val_loss: 0.3214
Epoch 9/25
285/285 [=====] - 113s 397ms/step - loss: 0.2392 - val_loss: 0.2772
Epoch 10/25
285/285 [=====] - 115s 403ms/step - loss: 0.2312 - val_loss: 0.2447
Epoch 11/25
285/285 [=====] - 112s 393ms/step - loss: 0.2206 - val_loss: 0.2448
Epoch 12/25
285/285 [=====] - 114s 399ms/step - loss: 0.2121 - val_loss: 0.2753
Epoch 13/25
285/285 [=====] - 113s 397ms/step - loss: 0.2039 - val_loss: 0.4078
Epoch 14/25
285/285 [=====] - 114s 399ms/step - loss: 0.1957 - val_loss: 0.2578
Epoch 15/25
285/285 [=====] - 113s 397ms/step - loss: 0.1914 - val_loss: 0.3132
Epoch 16/25
285/285 [=====] - 114s 399ms/step - loss: 0.1766 - val_loss: 0.2542
Epoch 17/25
285/285 [=====] - 116s 408ms/step - loss: 0.1747 - val_loss: 0.2660
Epoch 18/25
285/285 [=====] - 115s 403ms/step - loss: 0.1687 - val_loss: 0.2459
Epoch 19/25
285/285 [=====] - 114s 400ms/step - loss: 0.1628 - val_loss: 0.2630
Epoch 20/25
285/285 [=====] - 114s 399ms/step - loss: 0.1547 - val_loss: 0.3165
Epoch 21/25
285/285 [=====] - 115s 401ms/step - loss: 0.1486 - val_loss: 0.3829
Epoch 22/25
285/285 [=====] - 115s 403ms/step - loss: 0.1419 - val_loss: 0.2635
Epoch 23/25
285/285 [=====] - 114s 400ms/step - loss: 0.1328 - val_loss: 0.2537
Epoch 24/25
285/285 [=====] - 115s 405ms/step - loss: 0.1287 - val_loss: 0.2343
Epoch 25/25
285/285 [=====] - 114s 400ms/step - loss: 0.1232 - val_loss: 0.3024

```

Figure 10. DenseNet Run

optimizer, and regularization techniques (e.g., dropout rate) to find optimal settings that improve model performance. Data augmentation: Apply data augmentation techniques to expand the training dataset artificially. This can include random rotations, translations, flips, and zooms on the training images, which helps the model generalize better and reduces overfitting. Optimize training process: Experiment




Figure 11. CNN vs DenseNet Comparison

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_11 (Conv2D)	(None, 110, 110, 128)	9728
batch_normalization_11 (Batch Normalization)	(None, 110, 110, 128)	512
conv2d_12 (Conv2D)	(None, 108, 108, 256)	295168
batch_normalization_12 (Batch Normalization)	(None, 108, 108, 256)	1024
conv2d_13 (Conv2D)	(None, 106, 106, 256)	590080
batch_normalization_13 (Batch Normalization)	(None, 106, 106, 256)	1024
conv2d_14 (Conv2D)	(None, 104, 104, 256)	590080
batch_normalization_14 (Batch Normalization)	(None, 104, 104, 256)	1024
conv2d_15 (Conv2D)	(None, 102, 102, 128)	295040
batch_normalization_15 (Batch Normalization)	(None, 102, 102, 128)	512
conv2d_16 (Conv2D)	(None, 100, 100, 128)	147584
batch_normalization_16 (Batch Normalization)	(None, 100, 100, 128)	512
conv2d_17 (Conv2D)	(None, 98, 98, 128)	147584
batch_normalization_17 (Batch Normalization)	(None, 98, 98, 128)	512
conv2d_18 (Conv2D)	(None, 96, 96, 64)	73792
batch_normalization_18 (Batch Normalization)	(None, 96, 96, 64)	256
conv2d_19 (Conv2D)	(None, 94, 94, 64)	36928
batch_normalization_19 (Batch Normalization)	(None, 94, 94, 64)	256
conv2d_20 (Conv2D)	(None, 92, 92, 64)	36928
batch_normalization_20 (Batch Normalization)	(None, 92, 92, 64)	256
conv2d_21 (Conv2D)	(None, 90, 90, 32)	18464
batch_normalization_21 (Batch Normalization)	(None, 90, 90, 32)	128
global_average_pooling2d_1 (Global Average Pooling)	(None, 32)	0
dropout_1 (Dropout)	(None, 32)	0
dense_1 (Dense)	(None, 128)	4224
dropout_2 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 5)	645
Total params:		2,252,261
Trainable params:		2,249,253
Non-trainable params:		3,008

Figure 12. CNN Model Summary

Model: "sequential_3"

Layer (type)	Output Shape	Param #
densenet121 (Functional)	(None, 7, 7, 1024)	7037504
global_average_pooling2d_2 (Global Average Pooling)	(None, 1024)	0
dropout_3 (Dropout)	(None, 1024)	0
dense_2 (Dense)	(None, 5)	5125
Total params:		7,042,629
Trainable params:		6,958,981
Non-trainable params:		83,648

Figure 13. DenseNet Model Summary

with different optimization algorithms or adaptive learning rate techniques (e.g., Adam, RMSprop) to improve convergence and training speed.

5. Bibliography

References

- [1] Russell, S., & Norvig, P. (2020). *Artificial Intelligence: A Modern Approach* (4th ed.). Upper Saddle River, NJ: Pearson. Available at: <https://aima.cs.berkeley.edu/index.html>
- [2] Luger, G. F. (2005). *Artificial Intelligence: Structures and Strategies for Complex Problem Solving* (5th ed.). Addison Wesley.
- [3] Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep learning*. MIT Press. Available at: <https://www.deeplearningbook.org/>
- [4] Dive into Deep Learning. Available at: <http://d2l.ai/>
- [5] Géron, A. (2019). *Hands-on machine learning with Scikit-Learn, Keras and TensorFlow: concepts, tools, and techniques to build intelligent systems* (2nd ed.). O'Reilly.